

# CS349: Generative Adversarial Networks (GANs)

**Asif Ekbal**

Department of Computer Science and Engineering  
Indian Institute of Technology Patna

# Generative Adversarial Network (GAN)

- **Generative**

- Learn a generative model

- **Adversarial**

- Trained in an adversarial setting

- **Networks**


- Use Deep Neural Networks

# Why Generative Models?


- **We've only seen discriminative models so far**
  - Given an image  $\mathbf{X}$ , predict a label  $\mathbf{Y}$
  - Estimates  $\mathbf{P}(\mathbf{Y}|\mathbf{X})$
- **Discriminative models have several key limitations**
  - Can't model  $\mathbf{P}(\mathbf{X})$ , i.e. the probability of seeing a certain image
  - Thus, can't sample from  $\mathbf{P}(\mathbf{X})$ , i.e. **can't generate new images**
- **Generative models (in general) cope with all of above**
  - Can model  $\mathbf{P}(\mathbf{X})$
  - Can generate new images/any other data sample

# Magic of GANs...

Ground Truth



Adversarial




# Magic of GANs...

Which one is Computer generated?



# Magic of GANs...



# Adversarial Training

- **Important points**

- We can generate adversarial samples to fool a discriminative model
- We can use those adversarial samples to make models robust
- We then require more effort to generate adversarial samples
- Repeat this and we get better discriminative model

- **GANs extend that idea to generative models**

- **Generator**: generate fake samples, tries to fool the *Discriminator*
- **Discriminator**: tries to distinguish between *real and fake samples*
- Train them against each other
- Repeat this and we get better *Generator* and *Discriminator*

- So far we have looked at generative models which explicitly model the joint probability distribution or conditional probability distribution





- GANs take a different approach to this problem where the idea is to sample from a simple tractable distribution (say,  $z \sim N(0, I)$ ) and then learn a complex transformation from this to the training distribution



- GANs take a different approach to this problem where the idea is to sample from a simple tractable distribution (say,  $z \sim N(0, I)$ ) and then learn a complex transformation from this to the training distribution
- In other words, we will take a  $z \sim N(0, I)$ , learn to make a series of complex transformations on it so that the output looks as if it came from our training distribution

- What can we use for such a complex transformation?


- What can we use for such a complex transformation? **A Neural Network**

- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network?

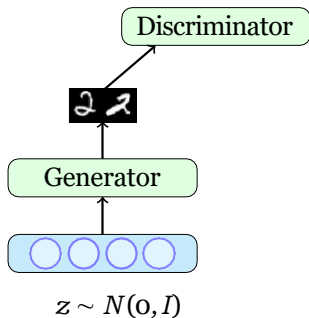
- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? **Using a two player game**

- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game:


- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game: a generator








- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game: a generator and **a discriminator**



- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game: a generator and a discriminator
- The job of the generator is to produce images which look so natural that the discriminator thinks that the images came from the real data distribution




- What can we use for such a complex transformation? A Neural Network
- How do you train such a neural network? Using a two player game
- There are two players in the game: a generator and a discriminator
- The job of the generator is to produce images which look so natural that the discriminator thinks that the images came from the real data distribution
- The job of the discriminator is to get better and better at distinguishing between true images and generated (fake) images


# View of GAN


- The simplest way of looking at a GAN is as a *generator network* that is trained to produce realistic samples by introducing an adversary i.e. the *discriminator network*, whose job is to detect if a given sample is “real” or “fake”
- Discriminator is a dynamically-updated evaluation metric for the tuning of the generator
- Both, the generator and discriminator continuously improve until an equilibrium point is reached
- **Generator**
  - improves as it receives feedback as to how well its generated samples managed to fool the discriminator
- **Discriminator**
  - improves by being shown not only the “fake” samples generated by the generator, but also “real” samples drawn from a real-life distribution
  - learns what generated samples look like and what real samples look like, thus enabling it to give better feedback to the generator

- So let's look at the full picture




- So let's look at the full picture
- Let  $G_\varphi$  be the generator and  $D_\theta$  be the discriminator ( $\varphi$  and  $\theta$  are the parameters of  $G$  and  $D$ , respectively)





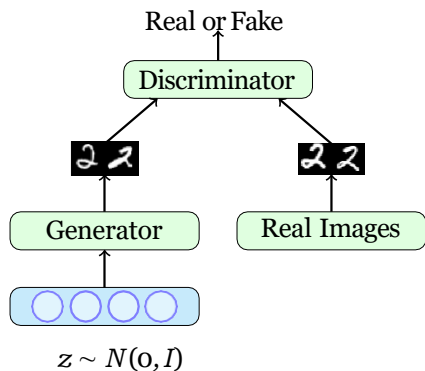
- So let's look at the full picture
- Let  $G_\varphi$  be the generator and  $D_\theta$  be the discriminator ( $\varphi$  and  $\theta$  are the parameters of  $G$  and  $D$ , respectively)
- We have a neural network based generator which takes as input a noise vector  $z \sim N(0, I)$  and produces  $G_\varphi(z) = X$




- So let's look at the full picture
- Let  $G_\varphi$  be the generator and  $D_\theta$  be the discriminator ( $\varphi$  and  $\theta$  are the parameters of  $G$  and  $D$ , respectively)
- We have a neural network based generator which takes as input a noise vector  $z \sim N(0, I)$  and produces  $G_\varphi(z) = X$
- We have a neural network based discriminator which could take as input a real  $X$  or a generated  $X = G_\varphi(z)$  and classify the input as real/fake




- What should be the objective function of the overall network?






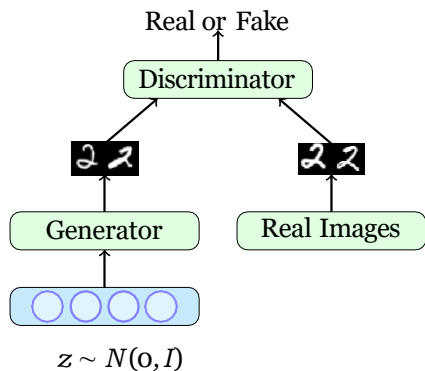
- What should be the objective function of the overall network?
- Let's look at the objective function of the generator first



- What should be the objective function of the overall network?
- Let's look at the objective function of the generator first
- Given an image generated by the generator as  $G_\phi(z)$  the discriminator assigns a score  $D_\theta(G_\phi(z))$  to it

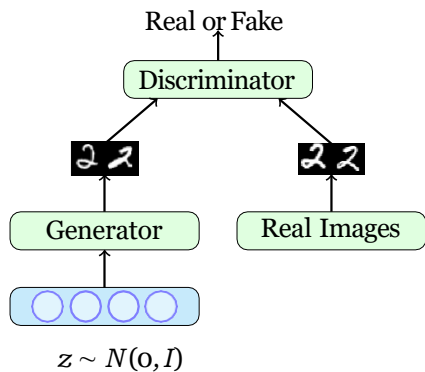



- What should be the objective function of the overall network?
- Let's look at the objective function of the generator first
- Given an image generated by the generator as  $G_\phi(z)$  the discriminator assigns a score  $D_\theta(G_\phi(z))$  to it
- This score will be between 0 and 1 and will tell us the probability of the image being real or fake



- What should be the objective function of the overall network?
- Let's look at the objective function of the generator first
- Given an image generated by the generator as  $G_\phi(z)$  the discriminator assigns a score  $D_\theta(G_\phi(z))$  to it
- This score will be between 0 and 1 and will tell us the probability of the image being real or fake
- For a given  $z$ , the generator would want to maximize  $\log D_\theta(G_\phi(z))$  (log likelihood) or minimize  $\log(1 - D_\theta(G_\phi(z)))$


- This is just for a single  $z$  and the generator would like to do this for all possible values of  $z$ ,





- This is just for a single  $z$  and the generator would like to do this for all possible values of  $z$ ,
- For example, if  $z$  was discrete and drawn from a uniform distribution (*i.e.*,  $p(z) = \frac{1}{N} \forall z$ ) then the generator's objective function would be

$$\min_{\phi} \sum_{i=1}^N \frac{1}{N} \log(1 - D_{\theta}(G_{\phi}(z)))$$



- This is just for a single  $z$  and the generator would like to do this for all possible values of  $z$ ,
- For example, if  $z$  was discrete and drawn from a uniform distribution (*i.e.*,  $p(z) = \frac{1}{N} \forall z$ ) then the generator's objective function would be

$$\min_{\phi} \sum_{i=1}^N \frac{1}{N} \log(1 - D_{\theta}(G_{\phi}(z)))$$


- However, in our case,  $z$  is continuous and not uniform ( $z \sim N(0, I)$ ) so the equivalent objective function would be

$$\min_{\phi} \int p(z) \log(1 - D_{\theta}(G_{\phi}(z)))$$


$$\min_{\phi} E_{z \sim p(z)} [\log(1 - D_{\theta}(G_{\phi}(z)))]$$




- Now let's look at the discriminator




- Now let's look at the discriminator
- The task of the discriminator is to assign a high score to real images and a low score to fake images






- Now let's look at the discriminator
- The task of the discriminator is to assign a high score to real images and a low score to fake images
- And it should do this for all possible real images and all possible fake images




- Now let's look at the discriminator
- The task of the discriminator is to assign a high score to real images and a low score to fake images
- And it should do this for all possible real images and all possible fake images
- In other words, it should try to maximize the following objective function

$$\max_{\theta} E_{x \sim p_{data}} [\log D_{\theta}(x)] + E_{z \sim p(z)} [\log(1 - D_{\theta}(G_{\phi}(z)))]$$

- If we put the objectives of the generator and discriminator together we get a minimax game



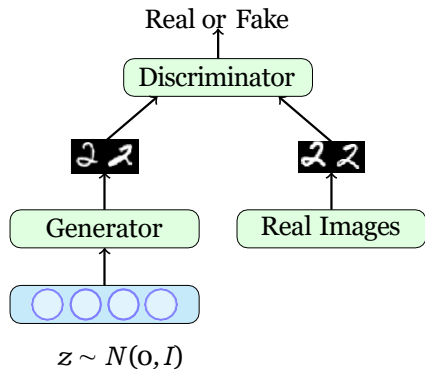
$$\min_{\varphi} \max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\varphi}(z)))]$$



- If we put the objectives of the generator and discriminator together we get a minimax game

$$\min_{\phi} \max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$


- The first term in the objective is only w.r.t. the parameters of the discriminator ( $\theta$ )



- If we put the objectives of the generator and discriminator together we get a minimax game

$$\min_{\varphi} \max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\varphi}(z)))]$$

- The first term in the objective is only w.r.t. the parameters of the discriminator ( $\theta$ )
- The second term in the objective is w.r.t. the parameters of the generator ( $\varphi$ ) as well as the discriminator ( $\theta$ )




- If we put the objectives of the generator and discriminator together we get a minimax game


$$\min_{\varphi} \max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\varphi}(z)))]$$

- The first term in the objective is only w.r.t. the parameters of the discriminator ( $\theta$ )
- The second term in the objective is w.r.t. the parameters of the generator ( $\varphi$ ) as well as the discriminator ( $\theta$ )
- The discriminator wants to maximize the second term whereas the generator wants to minimize it (hence it is a two-player game)



- So the overall training proceeds by alternating between these two step






- So the overall training proceeds by alternating between these two steps

- **Step 1:** Gradient Ascent on Discriminator

$$\max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$




- So the overall training proceeds by alternating between these two step

- **Step 1:** Gradient Ascent on Discriminator

$$\max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$

- **Step 2:** Gradient Descent on Generator

$$\min_{\phi} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))$$



- So the overall training proceeds by alternating between these two steps


- **Step 1:** Gradient Ascent on Discriminator

$$\max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$

- **Step 2:** Gradient Descent on Generator

$$\min_{\phi} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))$$

- In practice, the above generator objective does not work well and we use a slightly modified objective



- So the overall training proceeds by alternating between these two step

- **Step 1:** Gradient Ascent on Discriminator


$$\max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$


- **Step 2:** Gradient Descent on Generator

$$\min_{\phi} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))$$


- In practice, the above generator objective does not work well and we use a slightly modified objective
- Let us see why

- When the sample is likely fake, we want to give a feedback to the generator (using gradients)






- When the sample is likely fake, we want to give a feedback to the generator (using gradients)
- However, in this region where  $D(G(z))$  is close to 0, the curve of the loss function is very flat and the gradient would be close to 0



- When the sample is likely fake, we want to give a feedback to the generator (using gradients)
- However, in this region where  $D(G(z))$  is close to 0, the curve of the loss function is very flat and the gradient would be close to 0
- Trick: **Instead of minimizing the likelihood of the discriminator being correct, maximize the likelihood of the discriminator being wrong**





- When the sample is likely fake, we want to give a feedback to the generator (using gradients)
- However, in this region where  $D(G(z))$  is close to 0, the curve of the loss function is very flat and the gradient would be close to 0
- Trick: Instead of minimizing the likelihood of the discriminator being correct, maximize the likelihood of the discriminator being wrong
- In effect, the objective remains the same but the gradient signal becomes better

With that we are now ready to see the full algorithm for training GANs

---

---

1: procedure GAN TRAINING

2:     for number of training iterations do

3:         for k steps do

4:             • Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$

5:             • Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{data}(x)$

6:             • Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta} \frac{1}{m} \sum_{i=1}^m \left[ \log D_{\theta} \left( x^{(i)} \right) + \log \left( 1 - D_{\theta} \left( G_{\phi} \left( z^{(i)} \right) \right) \right) \right]$$

7:         end for

8:             • Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$

9:             • Update the generator by ascending its stochastic gradient

$$\nabla_{\phi} \frac{1}{m} \sum_{i=1}^m \left[ \log \left( D_{\theta} \left( G_{\phi} \left( z^{(i)} \right) \right) \right) \right]$$

10:     end for

11: end procedure

- We will now delve a bit deeper into the objective function used by GANs and see what it implies

- We will now delve a bit deeper into the objective function used by GANs and see what it implies
- Suppose we denote the true data distribution by  $p_{data}(x)$  and the distribution of the data generated by the model as  $p_G(x)$

- We will now delve a bit deeper into the objective function used by GANs and see what it implies
- Suppose we denote the true data distribution by  $p_{data}(x)$  and the distribution of the data generated by the model as  $p_G(x)$
- What do we wish should happen at the end of training?

- We will now delve a bit deeper into the objective function used by GANs and see what it implies
- Suppose we denote the true data distribution by  $p_{data}(x)$  and the distribution of the data generated by the model as  $p_G(x)$
- What do we wish should happen at the end of training?

$$p_G(x) = p_{data}(x)$$

- We will now delve a bit deeper into the objective function used by GANs and see what it implies
- Suppose we denote the true data distribution by  $p_{data}(x)$  and the distribution of the data generated by the model as  $p_G(x)$
- What do we wish should happen at the end of training?

$$p_G(x) = p_{data}(x)$$

- Can we prove this formally even though the model is not explicitly computing this density?

- We will now delve a bit deeper into the objective function used by GANs and see what it implies
- Suppose we denote the true data distribution by  $p_{data}(x)$  and the distribution of the data generated by the model as  $p_G(x)$
- What do we wish should happen at the end of training?

$$p_G(x) = p_{data}(x)$$

- Can we prove this formally even though the model is not explicitly computing this density?
- We will try to prove this over the next few slides



## Theorem

The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if and only if**  $p_G = p_{data}$

### Theorem

The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if and only if**  $p_G = p_{data}$

**is equivalent to**

### Theorem

The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if and only if**  $p_G = p_{data}$

**is equivalent to**

### Theorem

† **If**  $p_G = p_{data}$  then the global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **and**

### Theorem

The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if and only if**  $p_G = p_{data}$

**is equivalent to**

### Theorem

- 1 **If**  $p_G = p_{data}$  then the global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **and**
- 2 The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

## Outline of the Proof

**The ‘if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if**  $p_G = p_{data}$

**The ‘only if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

## Outline of the Proof

**The ‘if’ part:** The global minimum of the virtual training criterion

$C(G) = \max_D V(G, D)$  is achieved **if**  $p_G = p_{data}$

(a) Find the value of  $V(D, G)$  when the generator is optimal *i.e.*, when  $p_G = p_{data}$

**The ‘only if’ part:** The global minimum of the virtual training criterion

$C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

## Outline of the Proof

**The ‘if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if**  $p_G = p_{data}$

- (a) Find the value of  $V(D, G)$  when the generator is optimal *i.e.*, when  $p_G = p_{data}$
- (b) Find the value of  $V(D, G)$  for other values of the generator *i.e.*, for any  $p_G$  such that  $p_G \neq p_{data}$

**The ‘only if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

## Outline of the Proof

**The ‘if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if**  $p_G = p_{data}$

- (a) Find the value of  $V(D, G)$  when the generator is optimal *i.e.*, when  $p_G = p_{data}$
- (b) Find the value of  $V(D, G)$  for other values of the generator *i.e.*, for any  $p_G$  such that  $p_G \neq p_{data}$
- (c) Show that  $a < b \forall p_G \neq p_{data}$  (and hence the minimum  $V(D, G)$  is achieved when  $p_G = p_{data}$ )

**The ‘only if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

- Show that when  $V(D, G)$  is minimum then  $p_G = p_{data}$



## Outline of the Proof

**The ‘if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if**  $p_G = p_{data}$

- (a) Find the value of  $V(D, G)$  when the generator is optimal *i.e.*, when  $p_G = p_{data}$
- (b) Find the value of  $V(D, G)$  for other values of the generator *i.e.*, for any  $p_G$  such that  $p_G \neq p_{data}$
- (c) Show that  $a < b \forall p_G \neq p_{data}$  (and hence the minimum  $V(D, G)$  is achieved when  $p_G = p_{data}$ )

**The ‘only if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

- Show that when  $V(D, G)$  is minimum then  $p_G = p_{data}$

- First let us look at the objective function again

$$\min_{\phi} \max_{\theta} [\mathbb{E}_{x \sim p_{data}} \log D_{\theta}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta}(G_{\phi}(z)))]$$

- We will expand it to its integral form

$$\min_{\phi} \max_{\theta} \int_x p_{data}(x) \log D_{\theta}(x) + \int_z p(z) \log(1 - D_{\theta}(G_{\phi}(z)))$$

- Let  $p_G(X)$  denote the distribution of the  $X$ 's generated by the generator and since  $X$  is a function of  $z$  we can replace the second integral as shown below

$$\min_{\phi} \max_{\theta} \int_x p_{data}(x) \log D_{\theta}(x) + \int_x p_G(x) \log(1 - D_{\theta}(x))$$

- Okay, so our revised objective is given by

$$\min_{\phi} \max_{\theta} \int_x (p_{data}(x) \log D_{\theta}(x) + p_G(x) \log(1 - D_{\theta}(x))) dx$$

- Given a generator G, we are interested in finding the optimum discriminator D which will maximize the above objective function
- The above objective will be maximized when the quantity inside the integral is maximized  $\forall x$
- To find the optima we will take the derivative of the term inside the integral w.r.t.  $D$  and set it to zero

$$\frac{d}{d(D_{\theta}(x))} (p_{data}(x) \log D_{\theta}(x) + p_G(x) \log(1 - D_{\theta}(x))) = 0$$

$$p_{data}(x) \frac{1}{D_{\theta}(x)} + p_G(x) \frac{1}{1 - D_{\theta}(x)} (-1) = 0$$

$$\frac{p_{data}(x)}{D_{\theta}(x)} = \frac{p_G(x)}{1 - D_{\theta}(x)}$$

$$(p_{data}(x))(1 - D_{\theta}(x)) = (p_G(x))(D_{\theta}(x))$$

$$D_{\theta}(x) = \frac{p_{data}(x)}{p_G(x) + p_{data}(x)}$$

- This means for any given generator

$$D_G^*(G(x)) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$$

- Now the if part of the theorem says “if  $p_G = p_{data} \dots$ ”
- So let us substitute  $p_G = p_{data}$  into  $D_G^*(G(x))$  and see what happens to the loss functions

$$D_G^* = \frac{p_{data}}{p_{data} + p_G} = \frac{1}{2}$$

$$\begin{aligned} V(G, D_G^*) &= \int_x p_{data}(x) \log D(x) + p_G(x) \log (1 - D(x)) dx \\ &= \int_x p_{data}(x) \log \frac{1}{2} + p_G(x) \log \left(1 - \frac{1}{2}\right) dx \\ &= \log 2 \int_x p_G(x) dx - \log 2 \int_x p_{data}(x) dx \\ &= -2 \log 2 \quad = -\log 4 \end{aligned}$$

## Outline of the Proof

**The ‘if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if**  $p_G = p_{data}$

- (a) Find the value of  $V(D, G)$  when the generator is optimal *i.e.*, when  $p_G = p_{data}$
- (b) Find the value of  $V(D, G)$  for other values of the generator *i.e.*, for any  $p_G$  such that  $p_G \neq p_{data}$
- (c) Show that  $a < b \forall p_G \neq p_{data}$  (and hence the minimum  $V(D, G)$  is achieved when  $p_G = p_{data}$ )

**The ‘only if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

- Show that when  $V(D, G)$  is minimum then  $p_G = p_{data}$

- So what we have proved so far is that if the generator is optimal ( $p_G = p_{data}$ ) the discriminator's loss value is  $-\log 4$
- We still haven't proved that this is the minima
- For example, it is possible that for some  $p_G \neq p_{data}$ , the discriminator's loss value is lower than  $-\log 4$
- To show that the discriminator achieves its lowest value “if  $p_G = p_{data}$ ”, we need to show that for all other values of  $p_G$  the discriminator's loss value is greater than  $-\log 4$

- To show this we will get rid of the assumption that  $p_G = p_{data}$

$$\begin{aligned}
C(G) &= \int_x \left[ p_{data}(x) \log \left( \frac{p_{data}(x)}{p_G(x) + p_{data}(x)} \right) + p_G(x) \log \left( 1 - \frac{p_{data}(x)}{p_G(x) + p_{data}(x)} \right) \right] dx \\
&= \int_x \left[ p_{data}(x) \log \left( \frac{p_{data}(x)}{p_G(x) + p_{data}(x)} \right) + p_G(x) \log \left( \frac{p_G(x)}{p_G(x) + p_{data}(x)} \right) + (\log 2 - \log 2)(p_{data} + p_G) \right] dx \\
&= -\log 2 \int_x (p_G(x) + p_{data}(x)) dx \\
&\quad + \int_x \left[ p_{data}(x) \left( \log 2 + \log \left( \frac{p_{data}(x)}{p_G(x) + p_{data}(x)} \right) \right) + p_G(x) \left( \log 2 + \log \left( \frac{p_G(x)}{p_G(x) + p_{data}(x)} \right) \right) \right] dx \\
&= -\log 2(1 + 1) \\
&\quad + \int_x \left[ p_{data}(x) \log \left( \frac{p_{data}(x)}{\frac{p_G(x) + p_{data}(x)}{2}} \right) + p_G(x) \log \left( \frac{p_G(x)}{\frac{p_G(x) + p_{data}(x)}{2}} \right) \right] dx \\
&= -\log 4 + KL \left( p_{data} \parallel \frac{p_G(x) + p_{data}(x)}{2} \right) + KL \left( p_G \parallel \frac{p_G(x) + p_{data}(x)}{2} \right)
\end{aligned}$$

## Outline of the Proof

**The ‘if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **if**  $p_G = p_{data}$

- (a) Find the value of  $V(D, G)$  when the generator is optimal *i.e.*, when  $p_G = p_{data}$
- (b) Find the value of  $V(D, G)$  for other values of the generator *i.e.*, for any  $p_G$  such that  $p_G \neq p_{data}$
- (c) Show that  $a < b \forall p_G \neq p_{data}$  (and hence the minimum  $V(D, G)$  is achieved when  $p_G = p_{data}$ )

**The ‘only if’ part:** The global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

- Show that when  $V(D, G)$  is minimum then  $p_G = p_{data}$



- Okay, so we have

$$C(G) = -\log 4 + KL \left( p_{data} \parallel \frac{p_{data} + p_G}{2} \right) + KL \left( p_G \parallel \frac{p_{data} + p_G}{2} \right)$$

- We know that KL divergence is always  $\geq 0$

$$\therefore C(G) \geq -\log 4$$

- Hence the minimum possible value of  $C(G)$  is  $-\log 4$
- But this is the value that  $C(G)$  achieves when  $p_G = p_{data}$  (and this is exactly what we wanted to prove)
- We have, thus, proved the **if part** of the theorem

## Outline of the Proof

**The ‘if’ part:** The global minimum of the virtual training criterion

$C(G) = \max_D V(G, D)$  is achieved **if**  $p_G = p_{data}$

- (a) Find the value of  $V(D, G)$  when the generator is optimal *i.e.*, when  $p_G = p_{data}$
- (b) Find the value of  $V(D, G)$  for other values of the generator *i.e.*, for any  $p_G$  such that  $p_G \neq p_{data}$
- (c) Show that  $a < b \ \forall \ p_G \neq p_{data}$  (and hence the minimum  $V(D, G)$  is achieved when  $p_G = p_{data}$ )

**The ‘only if’ part:** The global minimum of the virtual training criterion

$C(G) = \max_D V(G, D)$  is achieved **only if**  $p_G = p_{data}$

- Show that when  $V(D, G)$  is minimum then  $p_G = p_{data}$

- Now let's look at the other part of the theorem

If the global minimum of the virtual training criterion  $C(G) = \max_D V(G, D)$  is achieved then

$$p_G = p_{data}$$

- We know that

$$C(G) = -\log 4 + KL \left( p_{data} \parallel \frac{p_{data} + p_g}{2} \right) + KL \left( p_G \parallel \frac{p_{data} + p_G}{2} \right)$$

- If the global minima is achieved then  $C(G) = -\log 4$  which implies that

$$KL \left( p_{data} \parallel \frac{p_{data} + p_g}{2} \right) + KL \left( p_G \parallel \frac{p_{data} + p_G}{2} \right) = 0$$

- This will happen only when  $p_G = p_{data}$  (you can prove this easily)
- In fact  $KL \left( p_{data} \parallel \frac{p_{data} + p_g}{2} \right) + KL \left( p_G \parallel \frac{p_{data} + p_G}{2} \right)$  is the Jensen-Shannon divergence between  $p_G$  and  $p_{data}$

$$KL \left( p_{data} \parallel \frac{p_{data} + p_g}{2} \right) + KL \left( p_G \parallel \frac{p_{data} + p_G}{2} \right) = JSD(p_{data} \parallel p_G)$$

which is minimum only when  $p_G = p_{data}$



