

Chapter 1: Introduction





Objectives

- Describe the **general organization of a computer system** and the role of interrupts
- Describe the **components in a modern**, multiprocessor computer system
- Illustrate the transition from **user mode to kernel mode**
- Discuss how operating systems are **used in various computing environments**
- Provide examples of **free and open-source operating systems**





What Does the Term Operating System Mean?

- An operating system is “fill in the blanks”
- What about:
 - Car
 - Airplane
 - Printer
 - Washing Machine
 - Toaster
 - Etc.





What is an Operating System?

- A program that acts as an **intermediary between a user of a computer and the computer hardware**
- Operating system goals:
 - **Execute user programs** and make solving user problems easier
 - Make the computer system **convenient to use**
 - Use the computer hardware in an ***efficient manner***






Computer System Structure

- Computer system can be divided into four components:
 - **Hardware** – provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - **Operating system**
 - ▶ **Controls and coordinates** use of hardware among various applications and users
 - **Application programs** – define the ways in which the system resources are used to solve the computing problems of the users
 - ▶ Word processors, compilers, web browsers, database systems, video games
 - **Users**
 - ▶ People, machines, other computers





Abstract View of Components of Computer





What Operating Systems Do

- Depends on the point of view
- Users want convenience, **ease of use** and **good performance**
 - Don't care about **resource utilization**
- But **shared computer** such as **mainframe** or **minicomputer** must keep all users happy
 - Operating system is a **resource allocator** and **control program** making **efficient use of HW** and managing execution of user programs
- Users of dedicate systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
- Mobile devices like **smartphones and tablets** are resource poor, optimized for usability and battery life
 - Mobile user interfaces such as touch screens, voice recognition
- Some computers **have little or no user interface**, such as embedded computers in devices and automobiles
 - Run primarily without user intervention





Defining Operating Systems

- Term OS covers many roles
 - Because of myriad designs and uses of OSES
 - Present in toasters through ships, spacecraft, game machines, TVs and industrial control systems
 - Born when **fixed use computers for military** became more **general purpose** and needed resource management and program control





Operating System Definition

- No universally accepted definition
- “Everything a vendor ships when you order an operating system” is a good approximation
 - But varies wildly
- “The one program running at all times on the computer” is the **kernel**, part of the operating system
- Everything else is either
 - A **system program** (ships with the operating system, but not part of the kernel) , or
 - An **application program**, all programs not associated with the operating system
- Today’s OSES for general purpose and mobile computing also include **middleware** – a set of software frameworks that provide **additional services to application developers** such as **databases, multimedia, graphics**






Overview of Computer System Structure





Computer System Organization

- Computer-system operation
 - One or more CPUs, device controllers connect through common **bus** providing access to shared memory
 - Concurrent execution of CPUs and devices **competing for memory cycles**





Computer-System Operation

- I/O devices and the CPU can **execute concurrently**
- Each **device controller** is in charge of a particular **device type**
- Each device controller has a **local buffer**
- Each device controller type has an operating system **device driver** to manage it
- CPU moves data from/to main memory to/from local buffers
- Device controller informs CPU that it has finished its operation by causing an **interrupt**






Common Functions of Interrupts

- An operating system is **interrupt driven**
- Interrupt transfers control to the **interrupt service routine (ISR)** generally, through the **interrupt vector**, which contains the **addresses of all the service routines**
- Interrupt architecture must save the **address of the interrupted instruction**
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request





Interrupt Timeline





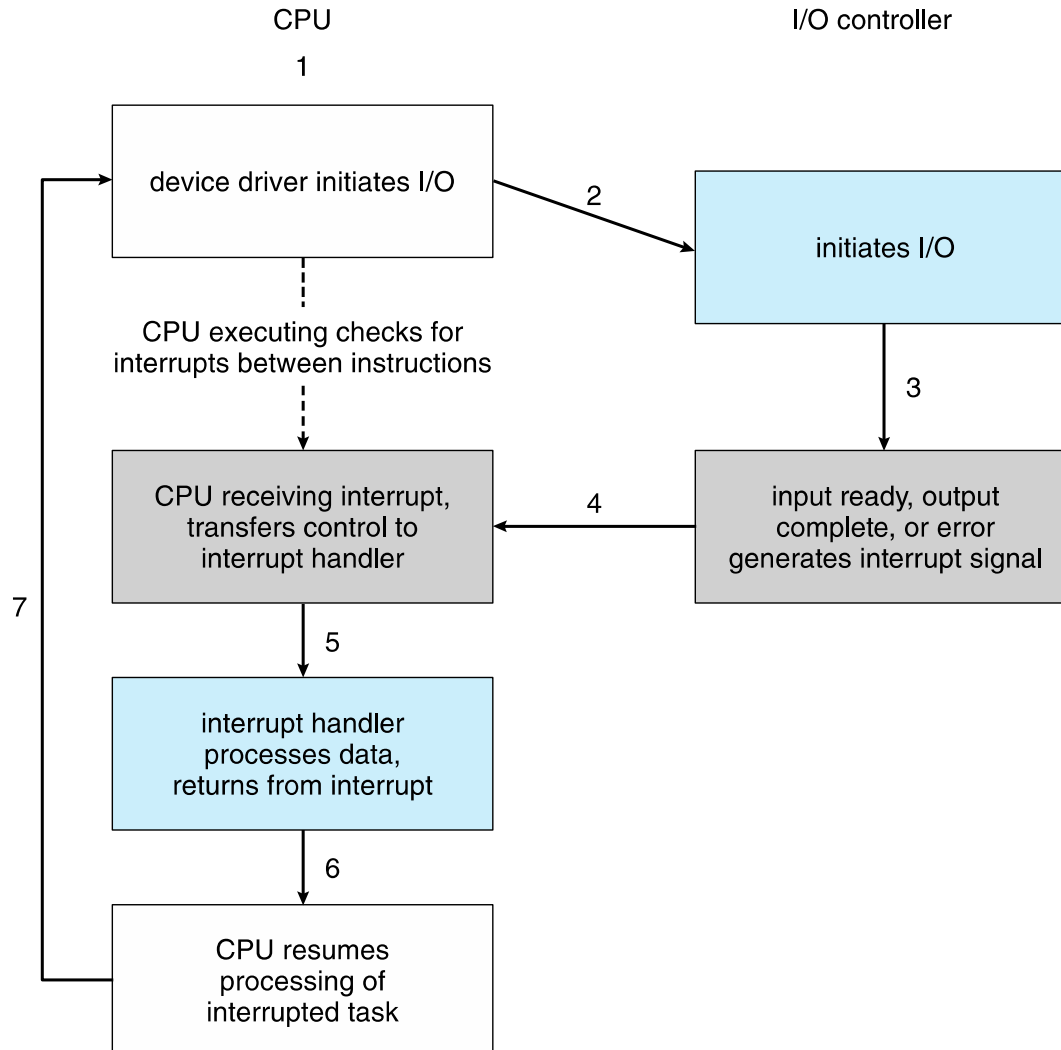
Interrupt Handling

- The operating system preserves the **state of the CPU** by **storing the registers** and **the program counter(PC)**
- Determines which type of interrupt has occurred:
- **Separate segments of code** determine **what action should be taken** for each type of interrupt





Interrupt-drive I/O Cycle





I/O Structure

- Two methods for handling I/O
 - After I/O starts, control returns to user program only upon I/O completion
 - After I/O starts, control returns to user program without waiting for I/O completion





I/O Structure (Cont.)

- After I/O starts, control returns to user program only upon I/O completion
 - Wait instruction **idles the CPU until the next interrupt**
 - Wait loop (contention for memory access)
 - **At most one I/O request is outstanding at a time**, no simultaneous I/O processing
- After I/O starts, control returns to user program without waiting for I/O completion
 - **System call** – request to the OS to allow user to wait for I/O completion
 - **Device-status table** contains entry for each I/O device indicating its **type, address, and state**
 - OS indexes into I/O device table to determine device status and to modify table entry to include interrupt





Storage Structure





Storage Structure

- Main memory – only large storage media that the CPU can access directly
 - **Random access**
 - Typically **volatile**
 - Typically **random-access memory** in the form of **Dynamic Random-access Memory (DRAM)**
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity





Storage Structure (Cont.)

- **Hard Disk Drives (HDD)** – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is **logically divided** into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the **logical interaction between the device and the computer**
- **Non-volatile memory (NVM)** devices– faster than hard disks, nonvolatile
 - Becoming more popular as capacity and performance increases, price drops






Storage Hierarchy

- Storage systems organized in hierarchy
 - Speed
 - Cost
 - Volatility
- **Caching** – copying information into faster storage system; **main memory can be viewed as a cache for secondary storage**






Storage-Device Hierarchy





How a Modern Computer Works



A von Neumann architecture





Direct Memory Access Structure

- Used for **high-speed I/O devices** able to transmit information at close to memory speeds
- Device controller **transfers blocks of data** from **buffer storage directly to main memory** without **CPU intervention**
- **Only one interrupt is generated per block**, rather than the one interrupt per byte





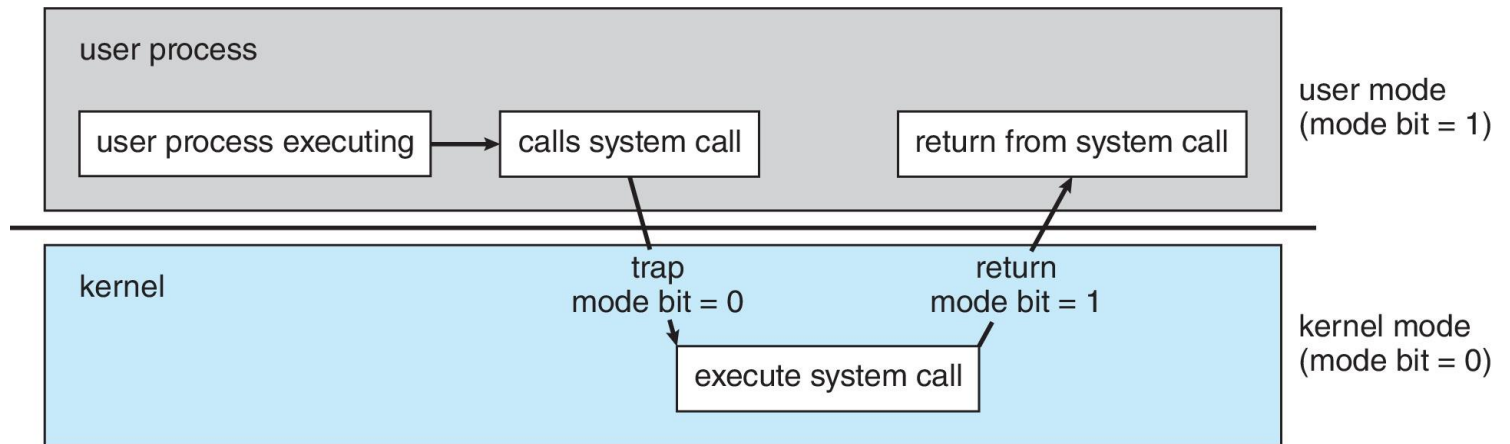
Dual-mode Operation

- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
- **Mode bit** provided by hardware
 - Provides ability to distinguish when **system is running user code or kernel code**.
 - When a user is running \Rightarrow **mode bit is “user”**
 - When kernel code is executing \Rightarrow **mode bit is “kernel”**
- How do we **guarantee that user does not explicitly set the mode bit to “kernel”**?
 - **System call** changes mode to kernel, return from call resets it to user
- Some instructions designated as **privileged**, only executable in kernel mode





Transition from User to Kernel Mode





Computer Startup

- **Bootstrap program** is loaded at power-up or reboot
 - Typically stored in ROM or EEPROM, generally known as **firmware**
 - Initializes all aspects of system (CPU registers, Device Controllers, etc.)
 - Loads **operating system kernel** and starts execution





Operating-System Operations

- Bootstrap program – **simple code to initialize the system**, load the kernel
- Kernel loads
- Starts **system daemons** (services provided outside of the kernel)
- Kernel **interrupt driven** (hardware and software)
 - Hardware interrupt by one of the devices
 - Software interrupt (**exception** or **trap**):
 - ▶ Software error (e.g., division by zero)
 - ▶ Request for operating system service – **system call**
 - ▶ Other process problems include infinite loop, *processes modifying each other or the operating system*





Multiprogramming (Batch system)

- Single user **cannot always keep CPU and I/O devices busy**
- Multiprogramming organizes jobs (**code and data**) so CPU always has one to execute
- A **subset of total jobs** in system is **kept in memory**
- One job selected and run via **job scheduling**
- When job has to wait (for I/O for example), OS switches to another job
- Completion time is **not very critical in batch processing**.
 - Objective is to **maximize the CPU utilization**.






Multitasking (Timesharing)

- A **logical extension of Batch systems**— the CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing
 - Objective is to **minimize the response time**.
 - **Response time** should be **< 1 second**
 - Each user has at least one program executing in memory \Rightarrow **process**
 - If several jobs **ready to run** at the same time \Rightarrow **CPU scheduling**
 - If processes don't fit in memory, **swapping** moves them in and out to run
 - **Virtual memory** allows execution of processes not completely in memory





Memory Layout for Multiprogrammed System





Timer

- Timer to **prevent infinite loop** (or process hogging resources)
 - Timer is set to **interrupt the computer after some time period**
 - Keep a counter that is **decremented by the physical clock**
 - Operating system **set the counter (privileged instruction)**
 - When counter zero generate an interrupt
 - Set up before scheduling process to **regain control or terminate program that exceeds allotted time**





Process Management

- A process is a program in execution. It is a unit of work within the system. Program is a *passive entity*; process is an *active entity*.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
- Process termination requires *reclaim of any reusable resources*
- *Single-threaded process* has one *program counter* specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- *Multi-threaded process* has *one program counter per thread*
- Typically system has many processes, some user, some operating system running concurrently on *one or more CPUs*
 - Concurrency by *multiplexing the CPUs* among the processes/threads





Process Management Activities

The **operating system is responsible** for the following activities in connection with process management:

- Creating and deleting both user and system processes
- Suspending and resuming processes
- Providing mechanisms for **process synchronization**
- Providing mechanisms for **process communication**
- Providing mechanisms for **deadlock handling**





Memory Management

- To **execute a program** all (or part) of the instructions must be in (primary) memory
- All (or part) of the **data that is needed** by the program must be in memory
- Memory management determines **what is in memory and when**
 - Optimizing CPU utilization and computer response to users
- Memory management activities
 - Keeping track of which parts of **memory are currently being used and by whom**
 - Deciding which processes (or parts thereof) and data to **move into and out of memory**
 - **Allocating and deallocating memory** space as needed





File-system Management

- OS provides uniform, logical view of information storage
 - Abstracts physical properties of storage devices to logical storage unit - **file**
 - Maps files onto physical media (magnetic disk, optical disk, etc)
- File-System management
 - Files usually organized into directories for ease of use.
 - Access control on most systems to determine who can access what
 - OS activities include
 - ▶ Creating and deleting files and directories
 - ▶ Primitives to manipulate files and directories
 - ▶ Mapping files onto secondary storage
 - ▶ Backup files onto stable (non-volatile) storage media





Mass-Storage Management

- Usually disks used to store data that does not fit in main memory or data that must be kept for a “long” period of time
- Proper management is of central importance
- Entire speed of computer operation hinges on disk subsystem and its algorithms
- OS activities
 - Mounting and unmounting
 - Free-space management
 - Storage allocation
 - Disk scheduling
 - Partitioning
 - Protection





Caching

- performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy





Protection and Security

- **Protection** – any mechanism for **controlling access of processes** or users to **resources defined by the OS**
- **Security** – **defense of the system** against **internal and external attacks**
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service
- Systems generally **first distinguish among users**, to determine who can do what
 - User identities (**user IDs**, security IDs) include name and associated number, **one per user**
 - **User ID then associated with all files**, processes of that user **to determine access control**
 - **Group identifier (group ID)** allows set of users to be defined and controls managed, then also associated with each process, file
 - **Privilege escalation** allows user to change to effective ID with more rights





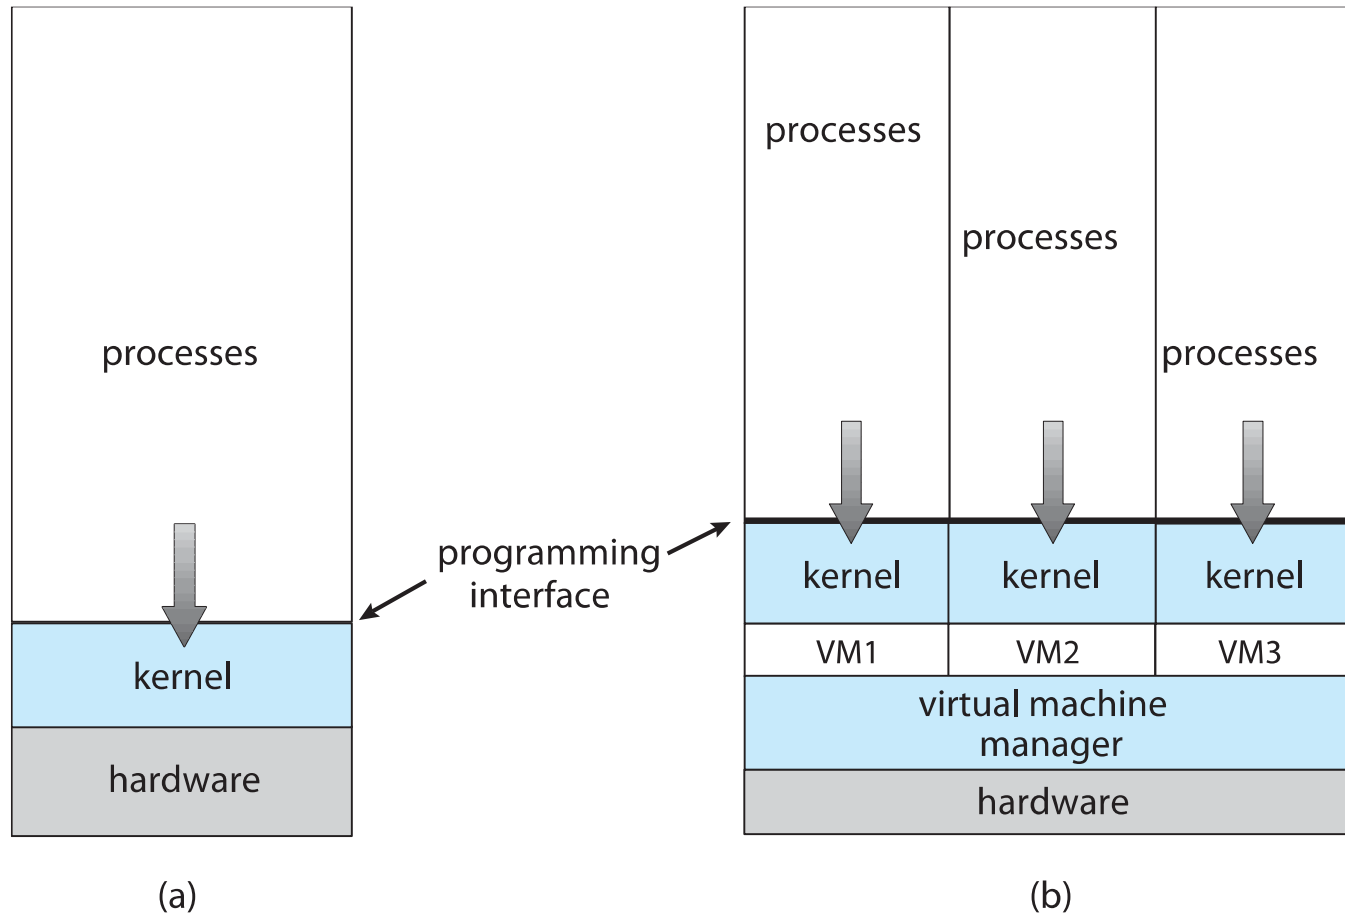
Virtualization

- **Emulation is** simulating **computer hardware in software**, is typically used when the **source CPU type is different from the target CPU type**.
 - Apple switched from the IBM Power CPU to the Intel x86 CPU for its desktop and laptop computers, it included an emulator **Rosetta**
- **Virtualization** is a technology that allows us to **abstract the hardware** of a single computer (the CPU, memory, disk drives, network interface cards etc) into **several different execution environments**.
- Creates the illusion that each separate environment (**virtual machine**) is running on its own private computer.
- Virtualization allows operating systems to **run as applications within other operating systems**.
- **Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - **VMM** (virtual machine Manager) provides virtualization services e.g., Vmware, Virtual Box





Computing Environments - Virtualization





Distributed Systems

- Collection of separate, possibly heterogeneous, systems networked together
 - **Network** is a communications path, **TCP/IP** most common
 - ▶ **Local Area Network (LAN)**
 - ▶ **Wide Area Network (WAN)**
 - ▶ **Metropolitan Area Network (MAN)**
 - ▶ **Personal Area Network (PAN)**
- **Network Operating System** provides features between systems across network
 - **Communication scheme** allows systems to exchange messages
 - Illusion of a single system





Computer System Architecture






Computer-System Architecture

- Most systems use a **single general-purpose processor**
 - Most systems have **special-purpose processors** as well
- **Multiprocessors** systems growing in use and importance
 - Also known as **parallel systems**, **tightly-coupled systems**
 - Advantages include:
 1. **Increased throughput**
 2. **Economy of scale**
 3. **Increased reliability** – graceful degradation or fault tolerance
 - Two types:
 1. **Asymmetric Multiprocessing** – each processor is assigned a specific task.
 2. **Symmetric Multiprocessing** – each processor performs all tasks






Symmetric Multiprocessing Architecture






Dual-Core Design


- The definition of multiprocessor has evolved over time and now **includes multicore systems**, in which multiple computing cores reside on a single chip i.e., Multi-chip and **multicore**
- more efficient than **multiple chips with single cores** because of faster on-chip communication
- multiple cores uses significantly **less power than multiple single-core** chips, hence suitable for mobile devices





Non-Uniform Memory Access System

- Adding additional CPUs to a multiprocessor system will increase computing; however, the concept does **not scale very well**
 - contention for the system bus **becomes a bottleneck**
 - **provide each CPU (or group of CPUs) with its own local memory**
- **Non-uniform memory access**
 - The CPUs are connected by a shared system interconnect, so that all CPUs share one physical address space

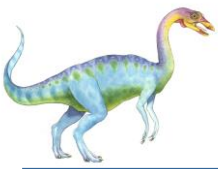





Clustered Systems

- Like multiprocessor systems, but multiple systems working together
- Clustered systems differ from the multiprocessor systems as they are **composed of two or more individual** (multicore) **systems**.
 - Usually sharing storage via a **storage-area network (SAN)**
 - Provides a **high-availability** service which survives failures
 - ▶ **Asymmetric clustering** has one machine in hot-standby mode to monitor the active server
 - ▶ **Symmetric clustering** has multiple nodes running applications, monitoring each other
 - Some clusters are for **high-performance computing (HPC)**
 - ▶ Applications must be written to use **parallelization**
 - Some have **distributed lock manager (DLM)** to avoid conflicting operations





Clustered Systems





Computer System Environments





Computing Environments

- Traditional
- Mobile
- Client Server
- Peer-to-Peer
- Cloud computing
- Real-time Embedded





Traditional

- Stand-alone general-purpose machines
- But blurred as most systems interconnect with others (i.e., the Internet)
- **Portals** provide web access to internal systems
- **Network computers (clients)** are like Web terminals
- Mobile computers interconnect via **wireless networks**
- Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks





Mobile


- Handheld smartphones, tablets, etc.
- What is the **functional difference between them** and a “traditional” laptop?
- Extra feature – more OS features (GPS, gyroscope)
- Allows new types of apps like ***augmented reality***
- Use IEEE 802.11 wireless, or cellular data networks for connectivity
- Leaders are **Apple iOS** and **Google Android**





Client Server


- Client-Server Computing
 - Dumb terminals supplanted by smart PCs
 - Many systems now **servers**, responding to requests generated by **clients**
 - ▶ **Compute-server system** provides an interface to client to request services (i.e., database)
 - ▶ **File-server system** provides interface for clients to store and retrieve files (i.e, web server)





Peer-to-Peer

- Another model of distributed system
- P2P does **not distinguish clients and servers**
 - Instead all nodes are considered peers
 - May each act as client, server or both
 - Node must join P2P network
 - ▶ Registers its service with **central lookup service on network**, or
 - ▶ Broadcast request for service and respond to requests for service via ***discovery protocol***
 - Examples include Torrent, Skype





Cloud Computing

- Delivers **computing, storage, even apps** as a service across a network
- Logical extension of virtualization because it uses virtualization as the base for its functionality.
 - Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, **pay based on usage**





Cloud Computing (Cont.)


- Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components
 - Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)
 - Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
 - Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use)





Cloud Computing (cont.)

- Cloud computing environments composed of traditional OSES, plus VMMs, plus cloud management tools
 - Internet connectivity requires security like firewalls
 - Load balancers spread traffic across multiple applications






Real-Time Embedded Systems

- Real-time embedded systems most prevalent form of computers
 - Vary considerable, special purpose, limited purpose OS, real-time OS
 - Found everywhere, from car engines and manufacturing robots to optical drives and microwave ovens.
 - Some have OSES, some perform tasks without an OS
 - ▶ application-specific integrated circuits (**ASICs**)
- Real-time OS has well-defined fixed time constraints
 - Processing **must** be done within constraint
 - Correct operation only if constraints met





Free and Open-Source Operating Systems

- Operating systems made available in source-code format rather than just binary **closed-source** and **proprietary**
- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**
 - Free software and open-source software are two different ideas championed by different groups of people
 - ▶ <https://www.gnu.org/philosophy/open-source-misses-the-point.en.html>
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)
 - Use to run guest operating systems for exploration






Kernel Data Structure






Kernel Data Structures


- Many similar to standard programming data structures
- ***Singly linked list***



- ***Doubly linked list***



- ***Circular linked list***






Kernel Data Structures

- **Binary search tree**

left \leq right


- Search performance is $O(n)$
- **Balanced binary search tree** is $O(\lg n)$





Kernel Data Structures

- **Hash function** can create a **hash map**



- **Bitmap** – string of n binary digits representing the status of n items
- Linux data structures defined in ***include*** files `<linux/list.h>`, `<linux/kfifo.h>`, `<linux/rbtree.h>`



End of Chapter 1

