

Chapter 3: Processes





Outline

- Process Concept
- Process Scheduling
- Operations on Processes
- Interprocess Communication
- IPC in Shared-Memory Systems
- IPC in Message-Passing Systems
- Examples of IPC Systems
- Communication in Client-Server Systems





Objectives

- Identify the **separate components of a process** and **illustrate how they are represented and scheduled** in an operating system.
- Describe how **processes are created and terminated** in an operating system using system calls.
- Describe and contrast **interprocess communication** using **shared memory and message passing**.
- Design programs that uses **pipes and POSIX shared memory** to perform **interprocess communication**.
- Describe client-server communication using sockets and remote procedure calls.
- Design kernel modules that interact with the Linux operating system.





Process Concept

- An operating system executes a **variety of programs that run as a process**.
- **Multiple processes** can execute concurrently, with the CPU (or CPUs) **multiplexed among them**.
- **Process** – a program in execution; process execution must progress in sequential fashion. **No parallel execution of instructions of a single process**
- Multiple parts
 - The **executable code**, also called **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory **dynamically allocated during run time**





Process Concept (Cont.)


- Program is **passive** entity stored on disk (**executable file**); process is **active**
 - Program becomes process when an executable file is loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc.
- One program can be **associated with several processes**
 - Consider multiple users **executing the same program**





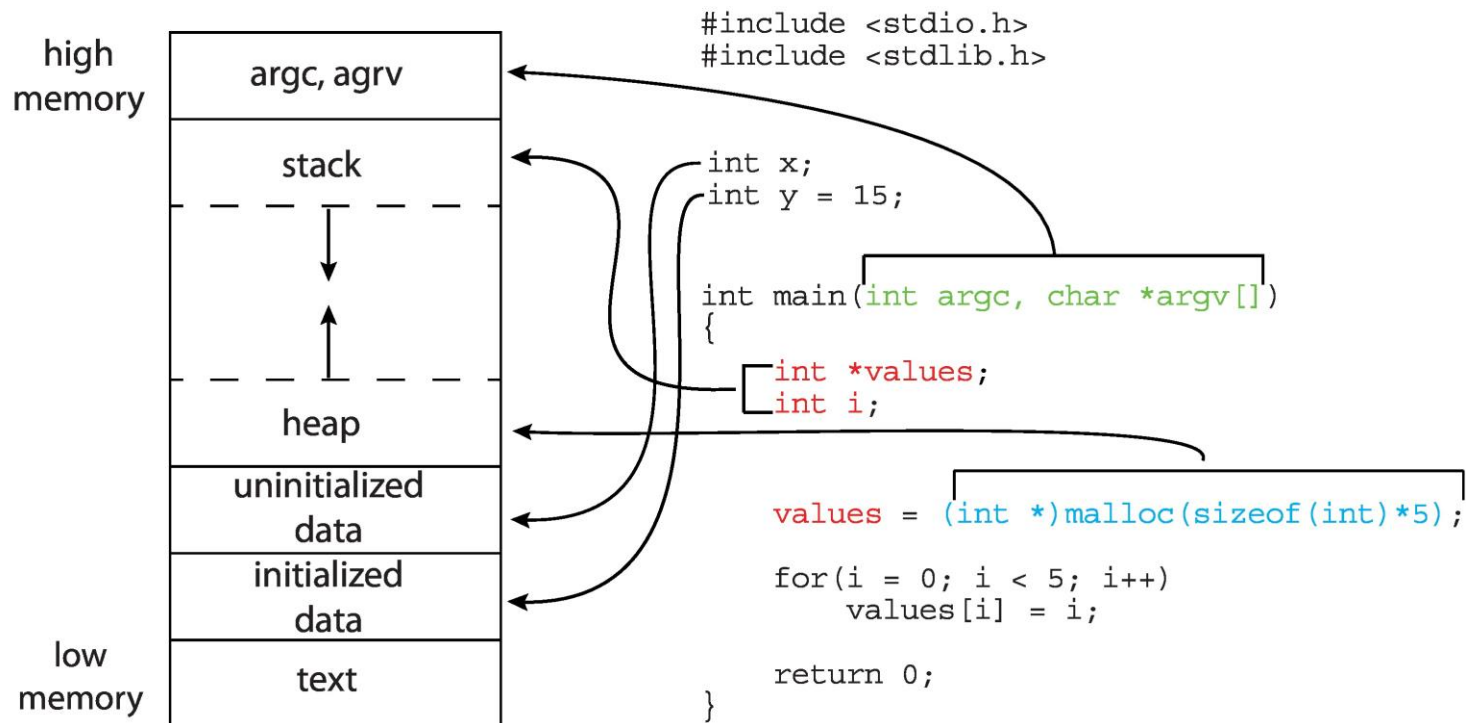
Process in Memory

- Sizes of the text and data sections are fixed (Why??)
- Operating system must ensure they do not **overlap** one another





Memory Layout of a C Program





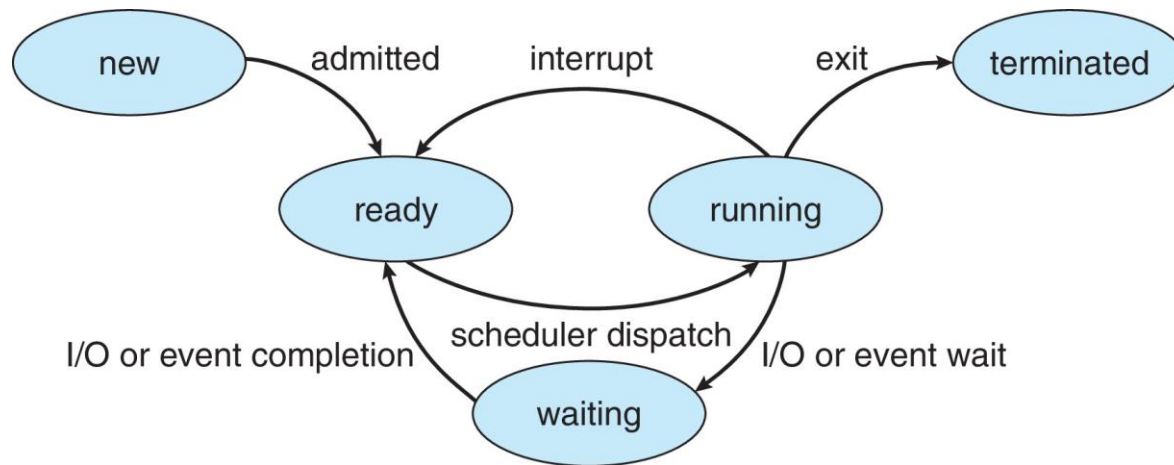
Process State

- As a process executes, it changes **state**
 - **New**: The process is being created
 - **Running**: Instructions are being executed
 - ▶ only one process can be running on any **processor core at any instant**
 - **Waiting**: The process is waiting for **some event to occur**
 - **Ready**: The process is **waiting to be assigned to a processor**
 - **Terminated**: The process has **finished execution**





Diagram of Process State





Process Control Block (PCB)

Information associated with each process(also called **task control block**)

- Process state – running, waiting, etc.
- Program counter – location of instruction to next execute
- CPU registers – contents of all process-centric registers e.g., accumulators, stack pointers, etc
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files

process state
process number
program counter
registers
memory limits
list of open files
...





Threads

- So far, process has a **single thread of execution**
- Consider having multiple program counters per process
 - Multiple locations can execute at once
 - ▶ Multiple threads of control -> **threads**
- PCB must have storage for thread details, multiple program counters
- Example: A multithreaded word processor could, for example, assign one thread to manage user input while another thread runs the spell checker.






Process Representation in Linux

Represented by the C structure `task_struct` in `<include/linux/sched.h>`

```
pid t_pid;                /* process identifier */
long state;               /* state of the process */
unsigned int time_slice   /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm;      /* address space of this
process */
```



- Within the Linux kernel, all active processes are represented using a doubly linked list of `task_struct`






Process Scheduling


- For a system with a single CPU core, there will never be more than one process running at a time.
- In case more processes than cores, excess processes will have to wait until a core is free.
- The number of processes currently in memory is known as the **degree of multiprogramming**.
- **Process scheduler** selects among available processes for next execution on CPU core
- Goal -- **Maximize CPU use**, quickly switch **processes onto CPU core**
- Maintains **scheduling queues** of processes
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Wait queues** – set of processes waiting for an event (i.e., I/O)
 - Processes migrate among the various queues (state diagram!!)






Ready and Wait Queues





Representation of Process Scheduling






Context Switch

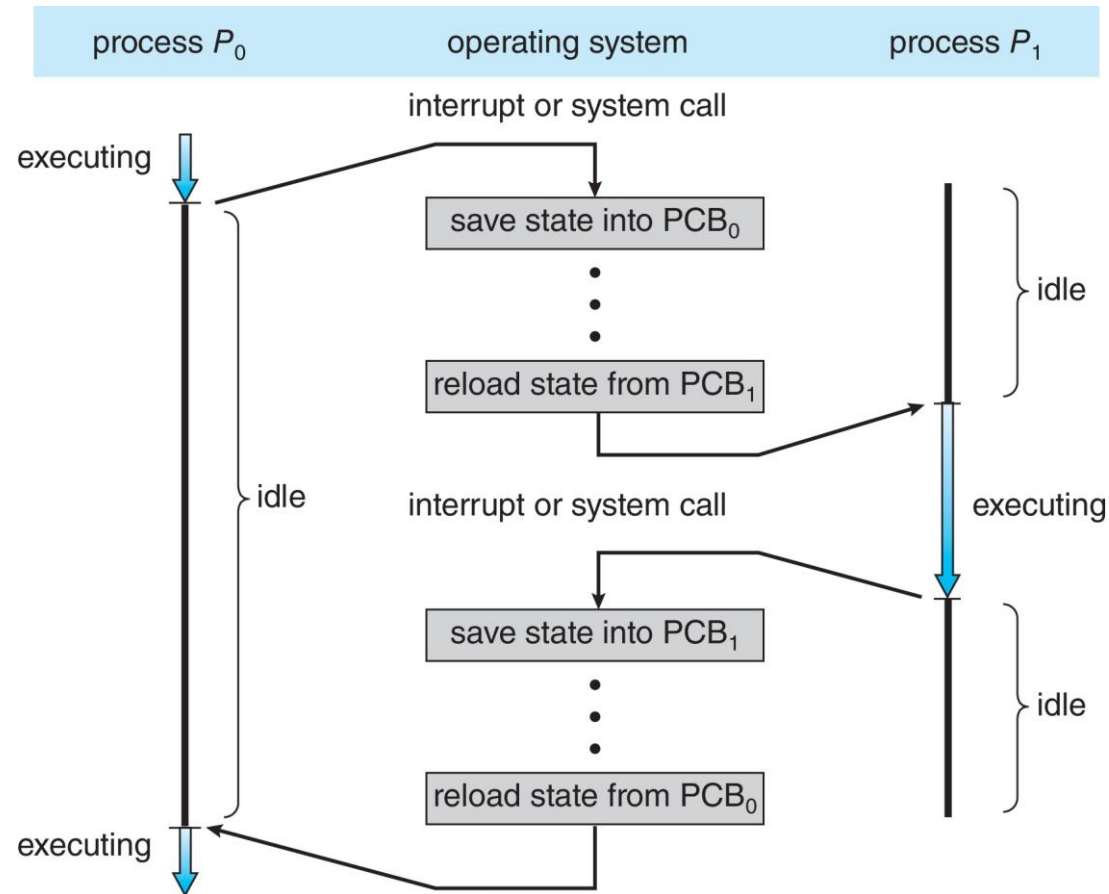
- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch **time is pure overhead**; the system does no useful work while switching
 - The more complex the OS and the PCB → **the longer the context switch**
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU → A context switch here simply requires changing the pointer to the current register set.





CPU Switch From Process to Process

A **context switch** occurs when the CPU switches from one process to another.





Multitasking in Mobile Systems

- Some mobile systems (e.g., early version of iOS) **allow only one process to run, others suspended**
- Due to limited screen size, user interface limits iOS provides for a
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes– in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs **foreground and background**, with fewer limits
 - Background **process uses a service to perform** tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use





Operations on Processes

- System must provide mechanisms for:
 - Process creation
 - Process termination





Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- **ps tree** command displays a tree of all processes in Linux
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate






Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program
 - Parent process calls **wait()** **waiting for the child to terminate**





A Tree of Processes in Linux





C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>


int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }


    return 0;
}
```






Process Creation (Cont.)

- After `fork()`, **one of the two processes** typically uses the `exec()` to replace the process's memory space with a new program.
- `execlp()` is a version of the `exec()` through which child process overlays its address space with the UNIX command `/bin/l`s
- The parent waits for the child process to complete with the `wait()`
 - When the child process completes, by either implicitly or explicitly, invoking `exit()` the parent process resumes from the call to `wait()`





Creating a Separate Process via Windows API

- Read by yourself





Process Termination

- Process executes last statement and then asks the operating system to delete it using the **exit()** system call.
 - Returns status data from child to parent (via **wait()**)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting, and the operating systems does not allow a child to continue if its parent terminates (**cascading termination**)
 - ▶ Normally initiated by the operating system





Process Termination

- The parent process may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```

- When a process terminates, its resources are deallocated
 - However, its entry in the **process table must remain there until the parent calls wait()**, because it contains the **process's exit status**
 - All processes transition to this **state when they terminate**
- If no parent waiting (did not invoke **wait()**) process is a **zombie**
- If parent terminated without invoking **wait()**, process is an **orphan**





Interprocess Communication

- Processes within a system may be *independent* or *cooperating*
- Cooperating process can **affect or be affected by other processes**, including sharing data
- Reasons for cooperating processes:
 - Information sharing (allow concurrent access to such information)
 - Computation speedup (we must break it into subtasks)
 - Modularity
- Cooperating processes need **interprocess communication (IPC)**





Interprocess Communication


- Two models of IPC
 - **Shared memory**
 - ▶ Processes can then exchange information by reading and writing data to the shared region
 - ▶ Shared memory can be faster than message passing (**Why??**)
 - **Message passing**
 - ▶ communication takes place by means of messages exchanged between the cooperating processes
 - ▶ Useful for exchanging smaller amounts of data, because no conflicts need be avoided
 - ▶ easier to implement in a distributed system






Communications Models

(a) Shared memory.



(a)

(b) Message passing.



(b)





IPC in Shared-Memory Systems

- Requires communicating processes to establish a **region of shared memory**
- Shared-memory region resides in the **address space of the process creating** the shared-memory segment.
 - OS tries to **prevent one process from accessing another process's** memory. Thus, processes agree **to remove this restriction**
- Exchange information by reading and writing data in the shared areas.
 - The form of the data and the location are determined by these processes and are **not under the OS's control**.





Producer-Consumer Problem

- Paradigm for cooperating processes:
 - *producer* process produces information that is consumed by a *consumer* process
 - ▶ e.g., a compiler may produce assembly code that is consumed by an assembler.
- Two variations:
 - **unbounded-buffer** places no practical limit on the size of the buffer:
 - ▶ Producer never waits
 - ▶ Consumer waits if there is no buffer to consume
 - **bounded-buffer** assumes that there is a fixed buffer size
 - ▶ Producer must wait if all buffers are full
 - ▶ Consumer waits if there is no buffer to consume





Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```





Producer Process – Shared Memory

```
item next_produced;

while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```





Consumer Process – Shared Memory

```
item next_consumed;

while (true) {
    while (in == out)
        ; /* empty buffer do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```

- Solution is correct, but can only use **BUFFER_SIZE-1** elements





What about Filling all the Buffers?

- Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers.
- We can do so by having an integer **counter** that keeps track of the number of full buffers.
- Initially, **counter** is set to 0.
- The integer **counter** is incremented by the producer after it produces a new buffer.
- The integer **counter** is and is decremented by the consumer after it consumes a buffer.





Producer

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE)  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```





Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```





Race Condition

- **counter++** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- **counter--** could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

- Consider this execution **interleaving** with “count = 5” initially:

S0: producer execute	register1 = counter	{register1 = 5}
S1: producer execute	register1 = register1 + 1	{register1 = 6}
S2: consumer execute	register2 = counter	{register2 = 5}
S3: consumer execute	register2 = register2 - 1	{register2 = 4}
S4: producer execute	counter = register1	{counter = 6}
S5: consumer execute	counter = register2	{counter = 4}





Race Condition (Cont.)

- Question - why was there no race condition in the first solution (where at most $N - 1$ buffers can be filled?)





IPC – Message Passing

- Processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(*message*)
 - **receive**(*message*)
- The *message* size is either fixed or variable





Message Passing (Cont.)

- If processes P and Q wish to communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via send/receive
- Implementation issues:
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between **every pair of communicating processes**?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?





Implementation of Communication Link

- Physical Implementation:
 - Shared memory
 - Hardware bus
 - Network
- Logical Implementation:
 - Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering (bounded or unbounded buffer)





Direct Communication

- Processes **must name each other explicitly**:
 - **send** (P , *message*) – send a message to process P
 - **receive**(Q , *message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically, processes need to know only each other's **identity to communicate**
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there **exists exactly one link**
 - The link may be unidirectional, but is usually bi-directional
- Two Variants
 - Both the **sender process and the receiver process must name the other to communicate** (***symmetry*** in addressing)
 - **Sender names the recipient**; the **recipient is not required to name the sender** (***asymmetry*** in addressing)





Indirect Communication

- Messages are **directed and received from mailboxes** (also referred to as ports)
 - Each mailbox has a **unique id**
 - Processes can communicate only if they **share a mailbox**
- Mailbox can be viewed **abstractly as an object** into which messages can be placed by processes and from which messages can be removed.
- Properties of communication link
 - Link established only if processes share a **common mailbox**
 - A link may be **associated with many processes**
 - Each pair of processes may **share several communication links**
 - Link may be unidirectional or bi-directional





Indirect Communication (Cont.)

- Operations
 - Create a new mailbox (port)
 - Send and receive messages through mailbox
 - Delete a mailbox
- Primitives are defined as:
 - **send**(*A, message*) – send a message to mailbox A
 - **receive**(*A, message*) – receive a message from mailbox A





Indirect Communication (Cont.)

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 , sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions depends on which of the following methods we choose
 - Allow a link to be associated with **at most two processes**
 - Allow only **one process at a time to execute a receive operation (How??)**
 - Allow the system to **select arbitrarily the receiver (may be round robin)**. Sender is notified who the receiver was.





Synchronization

Message passing may be either blocking or non-blocking

- **Blocking** is considered **synchronous**
 - **Blocking send** -- the sender is blocked until the message is received
 - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** -- the sender sends the message and continue
 - **Non-blocking receive** -- the receiver receives:
 - ▶ A valid message, or
 - ▶ Null message





Producer-Consumer: Message Passing

- Producer

```
message next_produced;  
while (true) {  
    /* produce an item in next_produced */  
  
    send(next_produced);  
}
```

- Consumer

```
message next_consumed;  
while (true) {  
    receive(next_consumed);  
  
    /* consume the item in next_consumed */  
}
```

- blocking send() and receive() statements





Buffering

- Whether communication is **direct or indirect**, messages exchanged by communicating processes reside in a **temporary queue**.
- Implemented in one of three ways
 1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length Sender never waits





Examples of IPC Systems – Windows

- Windows provides support for multiple operating environments, or **subsystems and application programs** communicate with these subsystems via a **message-passing mechanism**.
 - application programs are **clients** of a subsystem **server**
- Message-passing centric via **advanced local procedure call (ALPC)** facility
 - Only works between **processes on the same system**
 - Uses ports (like mailboxes) to establish and maintain communication channels






Examples of IPC Systems – Windows

- Communication works as follows:
 - The **client opens a handle** to the subsystem's **connection port** object.
 - The client sends a **connection request**.
 - The server creates **two private communication ports** and returns the handle to **one of them to the client**.
 - ▶ one for client–server messages, the other for server–client messages
 - The client and server use **the corresponding port handle to send messages** or callbacks and to listen for replies.
- For small messages (up to 256 bytes), the **port's message queue** is used as intermediate storage
- Larger messages must be passed through a **section object**





Local Procedure Calls in Windows





Pipes


- Acts as a channel allowing **two processes to communicate**
- First IPC mechanisms in early UNIX systems
- Issues:
 - Is communication unidirectional or bidirectional?
 - In the case of two-way communication, is it half or full-duplex?
 - Must there exist a relationship (i.e., **parent-child**) between the communicating processes?
 - Can the pipes be used over a network?
- **Ordinary pipes** – **cannot be accessed from outside the process that created it**. Typically, a parent process creates a pipe and uses it to communicate with a child process that it created.
- **Named pipes** – can be accessed without a parent-child relationship.





Ordinary Pipes

- Ordinary Pipes allow communication in standard **producer-consumer style**
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are **therefore unidirectional**
 - If two-way communication is required, two pipes must be used
- Require parent-child relationship between communicating processes



- Windows calls these **anonymous pipes**
- UNIX treats a pipe as a **special type of file**. Thus, pipes can be accessed using ordinary **read()** and **write()** system calls





Named Pipes

- Named Pipes are more powerful than ordinary pipes
- **Communication is bidirectional**
- **No parent-child relationship** is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems





Communications in Client-Server Systems

- Sockets
- Remote Procedure Calls






Sockets

- A **socket** is defined as an **endpoint for communication**
- **Concatenation of IP address and port** – a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are **well known**, used for standard services
 - FTP server listens to port 21; and a web, or HTTP, server listens to port 80
- Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running





Socket Communication





Sockets in Java

- Three types of sockets
 - **Connection-oriented (TCP)**
 - **Connectionless (UDP)**
 - **MulticastSocket** class— data can be sent to multiple recipients
- Consider this “Date” server in Java:

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





Sockets in Java

The equivalent Date client

```
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args) {
        try {
            /* make connection to server socket */
            Socket sock = new Socket("127.0.0.1",6013);

            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

            /* read the date from the socket */
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            /* close the socket connection*/
            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```





Remote Procedure Calls

- Read by yourself



End of Chapter 3

