

# Chapter 9: Main Memory

---





# Chapter 9: Memory Management

---

- Background
- Contiguous Memory Allocation
- Paging
- Structure of the Page Table
- Swapping
- Example: The Intel 32 and 64-bit Architectures
- Example: ARMv8 Architecture





# Objectives

---

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques,
- To provide a detailed description of the Intel Pentium, which supports both **pure segmentation and segmentation with paging**





# Background

---

- Program must be brought (from disk) into **memory and placed within a process for it to be run**
- The CPU fetches instructions from memory according to the PC.
  - Fetched instructions may cause **additional loading from and storing** to specific memory addresses.
- Memory unit only sees a stream of:
  - addresses + read requests, or
  - address + data + write requests
  - does not know how addresses are being generated (e.g., instruction counter, indexing, indirection, etc)
  - does not know the address is for **instructions or data**





# Background

---

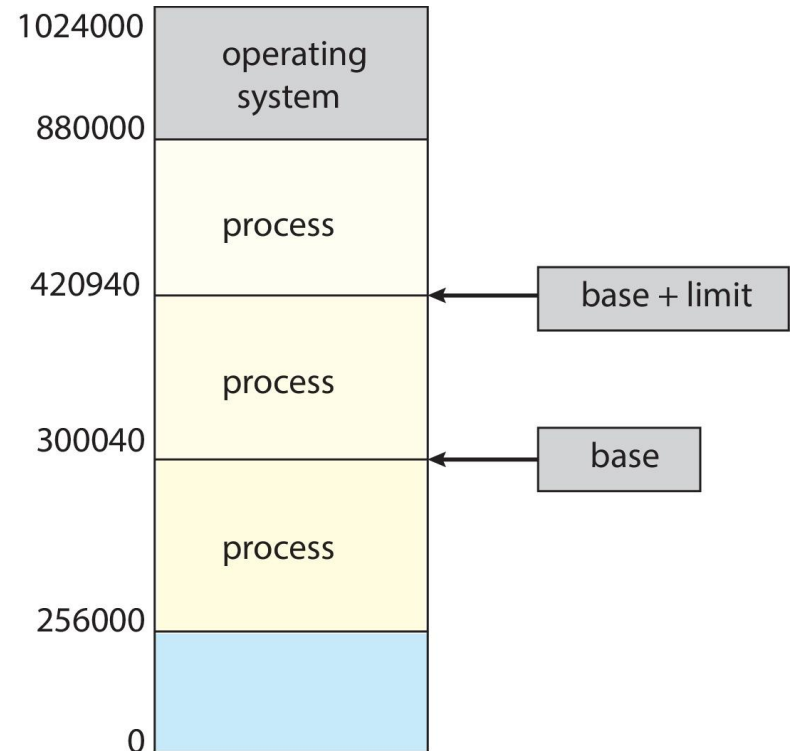
- Main memory and registers are only storage CPU can access directly
  - If the desired data are not in memory, they must be moved there before the CPU can operate on them
- Register are built into each CPU core.
  - Accessible within one cycle of the CPU clock.
- Main memory can take many cycles; processor needs to stall
- Cache sits between main memory and CPU registers





# Protection

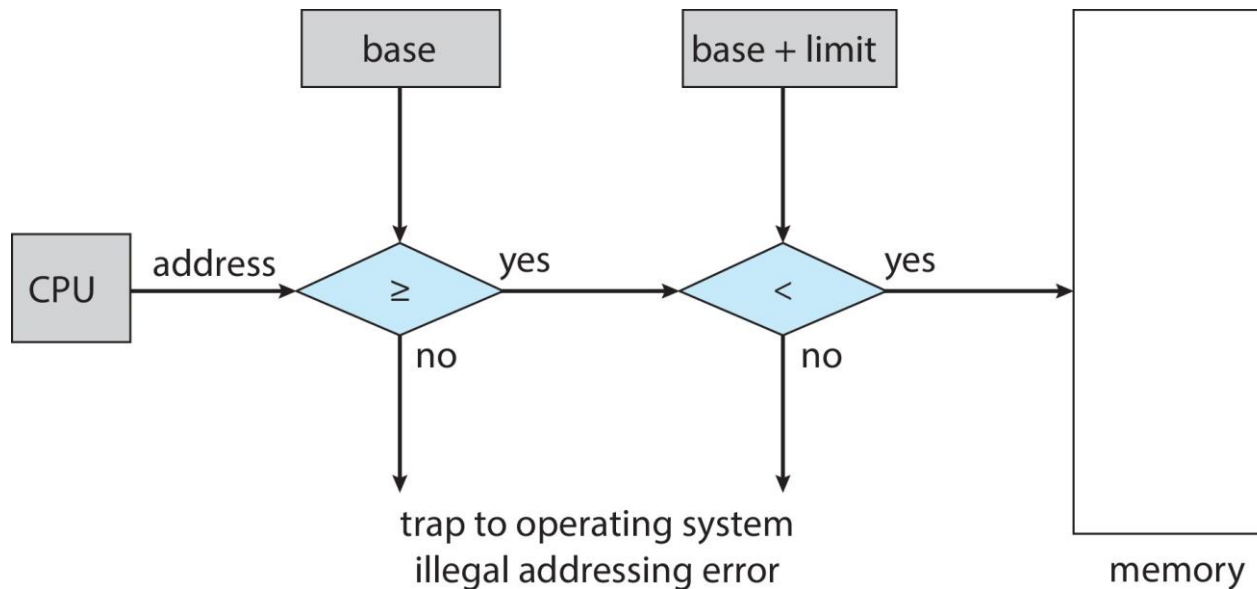
- Along with **relative speed of memory access** correct operation must be ensured.
- **Protection of memory** required to ensure correct operation
- Need to ensure that a process can access only those **addresses in its address space**.
  - Each process should have separate memory space
- We can provide this protection by using a pair of **base** and **limit registers** define the logical address space of a process





# Hardware Address Protection

- Protection of memory space is accomplished by having the CPU hardware compare every **address generated in user mode** with the **base and limit registers**.



- the instructions to loading the **base and limit registers** are privileged





# Address Binding

---

- Addresses represented in different ways at different stages of a program's life
  - Source code addresses usually symbolic
  - Compiled code addresses **bind** to relocatable addresses
    - ▶ i.e., “14 bytes from beginning of this module”
  - Linker or loader will bind relocatable addresses to absolute addresses
    - ▶ i.e., 74014
  - Each binding maps one address space to another







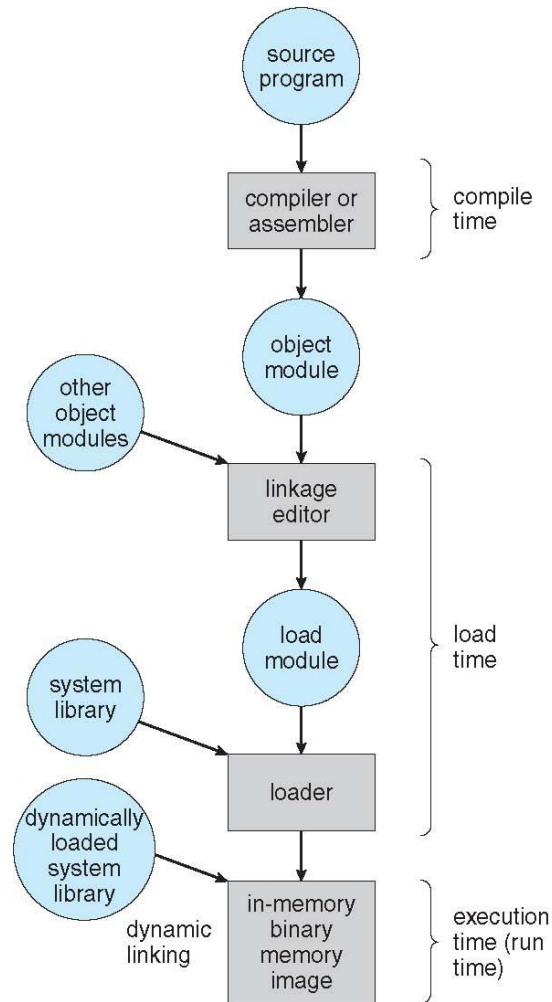
# Binding of Instructions and Data to Memory


- Address binding of instructions and data to memory addresses can happen at three different stages
  - **Compile time:** If memory location known a priori, **absolute code** can be generated; **must recompile code if starting location changes**
  - **Load time:** Must generate **relocatable code** if memory location is **not known at compile time**
  - **Execution time:** Binding delayed until run time **if the process can be moved during its execution** from one memory segment to another
    - ▶ Need hardware support for address maps (e.g., base and limit registers)





# Multistep Processing of a User Program





# Logical vs. Physical Address Space

---

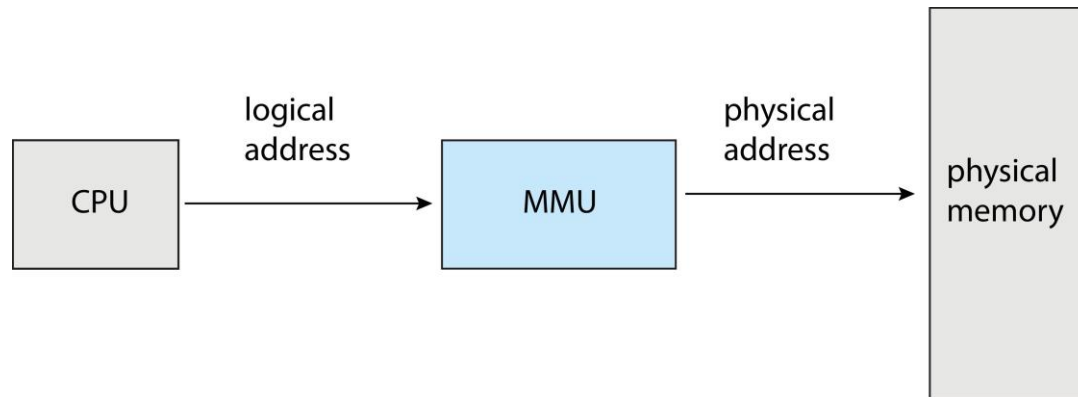
- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU; also referred to as **virtual address**
  - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the **same in compile-time and load-time address-binding schemes**; logical (virtual) and physical addresses **differ in execution-time** address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program





# Memory-Management Unit (MMU)

- Hardware device that at run time maps virtual to physical address




- Many methods possible, covered in the rest of this chapter





# Memory-Management Unit (Cont.)

- Consider **simple scheme**, which is a generalization of the base-register scheme.
- The base register now called **relocation register**
- The value in the relocation register is added to every address generated by a user process at the time it is sent to memory





# Dynamic Loading

- Routine is **not loaded until it is called**
- Better memory-space utilization; unused routine is **never loaded**
- All routines kept on disk in **relocatable load format**
- Useful when **large amounts of code** are needed to **handle infrequently occurring cases**
- When a routine needs to call another routine, the **calling routine first checks** the availability of other routine in the memory
  - If not, relocatable linking loader is called to load the desired routine and to update the **program's address tables** to reflect this change.
- No special support from the operating system is required
  - Implemented through program design (**How??**)
  - OS may help by providing libraries to implement dynamic loading





# Dynamic Linking

- **Static linking**
  - System library modules get fully included in executable modules
  - wasting both disk space and main memory usage
- **Dynamic linking** –linking postponed until execution time
- Small piece of code (i.e., **stub**), used as reference to the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Dynamic linking is particularly useful for system libraries
- Dynamically linked libraries can be shared among multiple processes, thus, also known as **shared libraries**
- Consider applicability to patching system libraries
  - Versioning may be needed





# Contiguous Allocation

---

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two partitions:
- We can place the operating system in either low memory addresses or high memory addresses. Linux and Windows place the operating system in high memory
  - User processes then held in low/high memory accordingly
  - Each process contained in single contiguous section of memory







# Contiguous Allocation (Cont.)

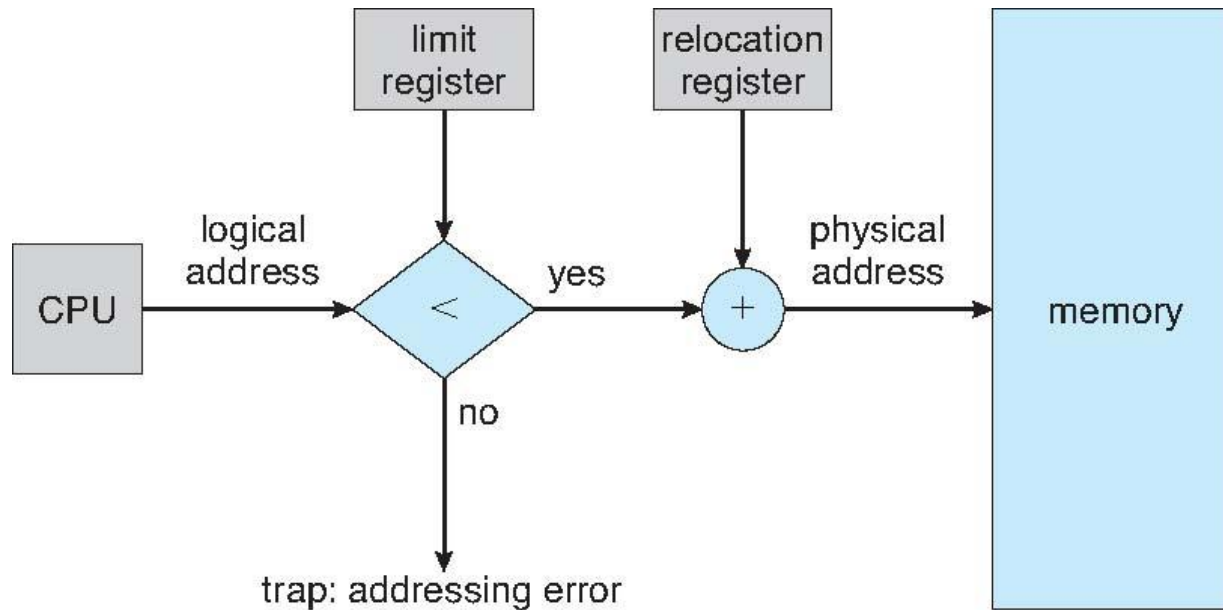
---

- **Relocation registers** used to protect user processes from each other, and from changing operating-system code and data
  - Relocation register contains value of **smallest physical address**
  - Limit register contains **range of logical addresses** – each logical address must be less than the limit register
  - MMU maps logical address *dynamically* by adding the value in the relocation register
- When the CPU scheduler selects a process for execution, the dispatcher loads the **relocation and limit registers** with the correct values.






# Hardware Support for Relocation and Limit Registers





# Variable Partition

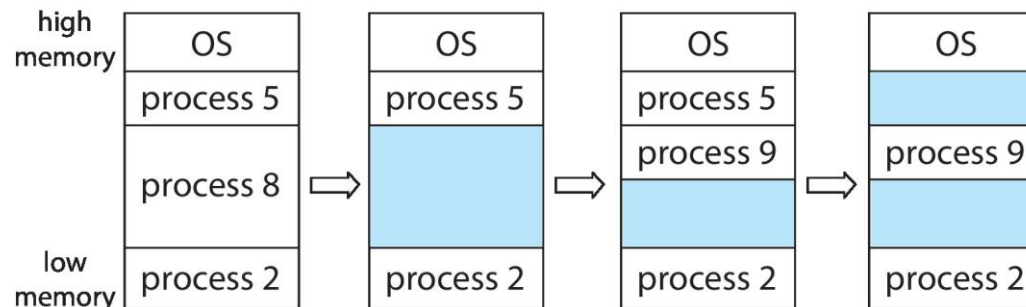
- OS keeps a table indicating which parts of memory are available and which are occupied.
  - Initially, one large block (**single hole**) of memory is available for user processes.






# Variable Partition

- Multiple-partition allocation
  - Degree of multiprogramming limited by **number of partitions**
  - **Variable-partition** sizes for efficiency (sized to a given **process' needs**)
  - **Hole** – block of available memory; holes of various size are scattered throughout memory
  - When a process arrives, it is allocated memory from a **hole large enough to accommodate it**
  - Process exiting frees its partition, **adjacent free partitions combined**
  - Operating system maintains information about:
    - a) **allocated partitions**
    - b) **free partitions (hole)**





# Dynamic Storage-Allocation Problem

---

How to satisfy a request of size  $n$  from a list of free holes?

- **First-fit:** Allocate the ***first*** hole that is big enough
  - Searching can start either at the beginning of the set of holes or at the location where the previous first-fit search ended.
- **Best-fit:** Allocate the ***smallest*** hole that is big enough; must search entire list, unless ordered by size
  - Produces the smallest leftover hole
- **Worst-fit:** Allocate the ***largest*** hole; must also search entire list
  - Produces the largest leftover hole

Based on simulation, **first-fit and best-fit** can be considered better than worst-fit in terms of speed and storage utilization





# Fragmentation

---

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Internal Fragmentation** – **allocated memory** may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- First fit statistical analysis reveals that given  $N$  blocks allocated,  $0.5 N$  blocks lost to fragmentation
  - **50-percent rule**





# Fragmentation (Cont.)

---

- Reduce external fragmentation by **compaction**
  - Shuffle memory contents to place all free memory together in one large block
  - Compaction is possible only if relocation is dynamic, and is done at execution time





# Paging

---

- Physical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available
  - Avoids **external fragmentation and the associated need for compaction**
- Divide physical memory into **fixed-sized blocks** called **frames**
  - Size is power of 2, between **512 bytes and 16 Mbytes**
- Divide programs **logical memory** into blocks of same size called **pages**
- Keep track of **all free frames**
  - Any page ( from any process ) can be placed into any available frame.
- Set up a **page table** to **translate logical to physical addresses**
- Still have **Internal fragmentation**








# Address Translation Scheme

- Address generated by CPU is divided into:
  - **Page number** ( $p$ ) – used as an index into a **page table** which contains **base address of each page** in physical memory
  - **Page offset** ( $d$ ) – combined with base address to define the **physical memory address** that is sent to the memory unit

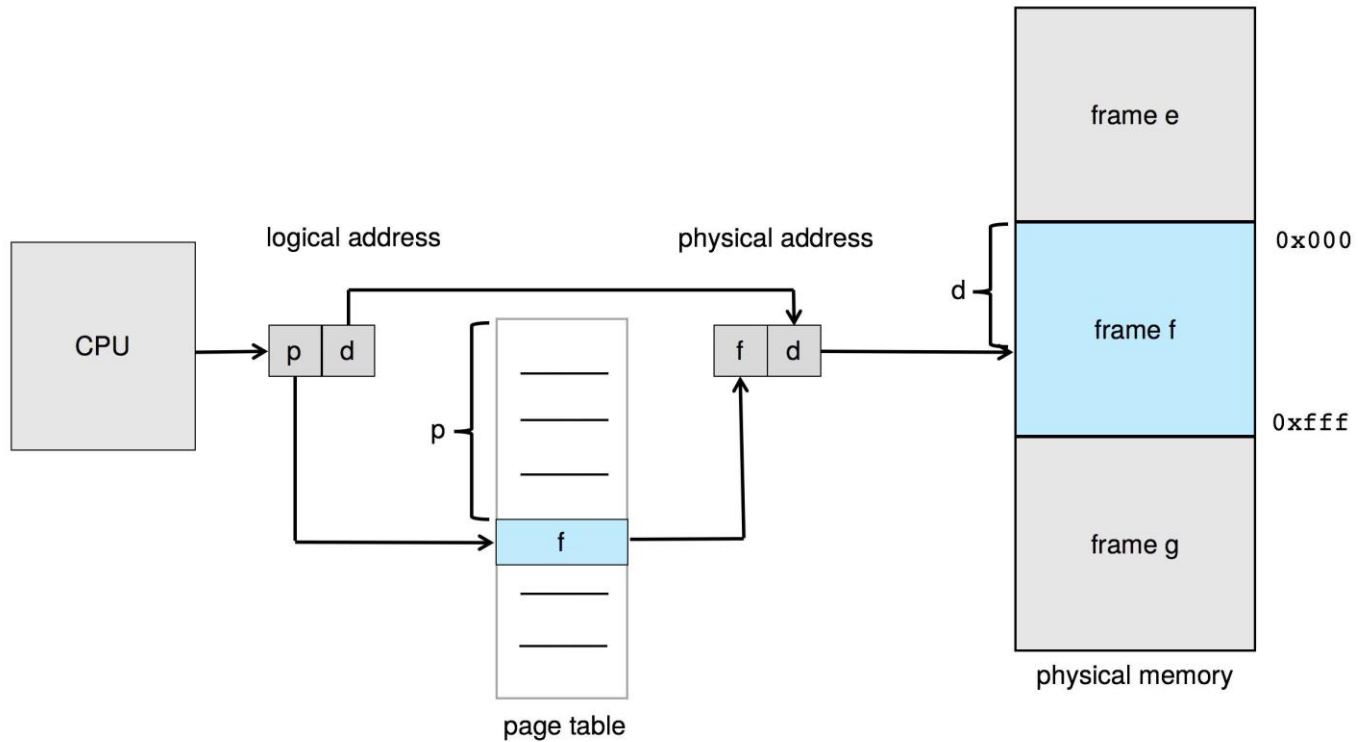


- For given logical address space  $2^m$  and page size  $2^n$





# Paging Hardware

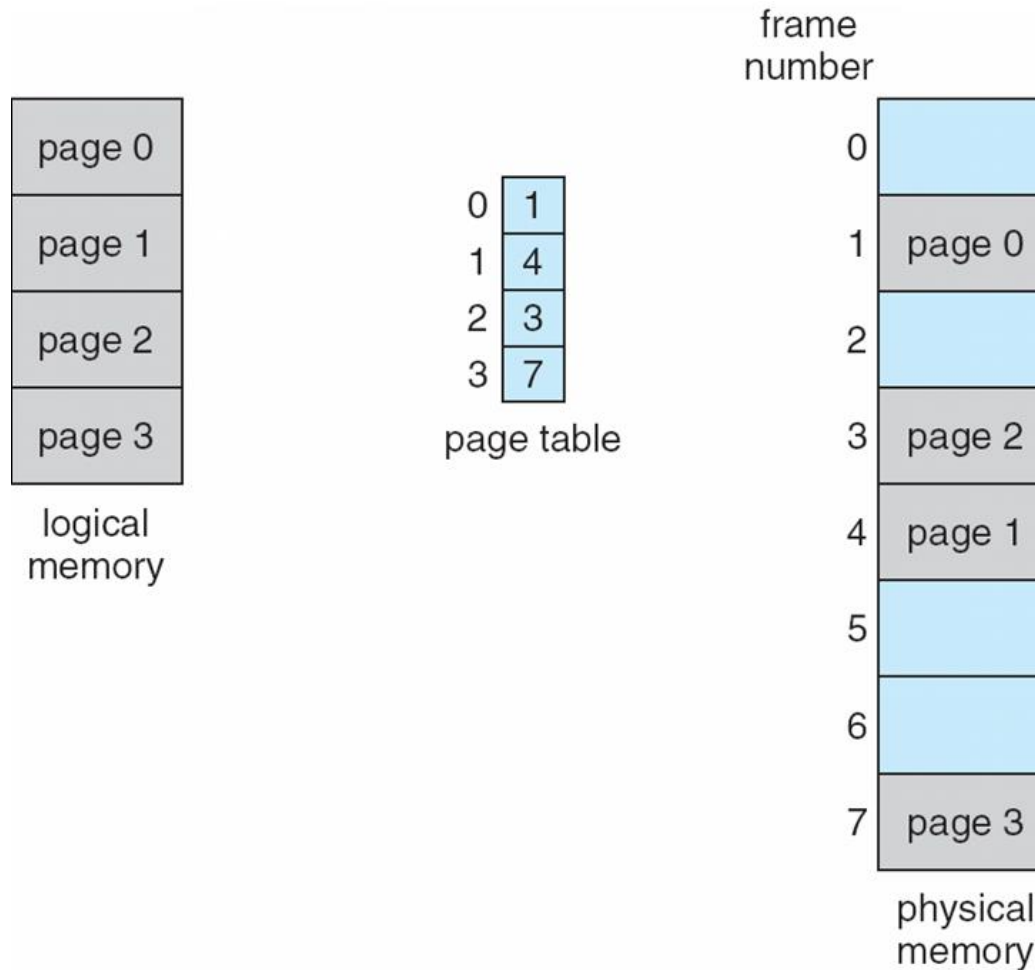


Note: Pages and frames are of same size.





# Paging Model of Logical and Physical Memory





# Paging Example

- Consider the following micro example, in which a process has 16 bytes of logical memory, mapped in 4 byte pages into 32 bytes of physical memory.

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory






# Paging -- Calculating internal fragmentation

- Page size = 2,048 bytes
- Process size = 72,766 bytes
- 35 pages + 1,086 bytes
- Internal fragmentation of  $2,048 - 1,086 = 962$  bytes
- Worst case fragmentation = 1 frame – 1 byte
- On average fragmentation = 1 / 2 frame size
- So small frame sizes desirable?
  - But each page table entry takes memory to track
  - Inefficient from the perspective of disk I/O
- With the advancement in processor page sizes are growing
  - On x86-64 systems, Windows 10 supports page sizes of 4 KB and 2 MB
- For Linux default page size (typically 4 KB) and an architecture dependent larger page size
  - System call `getpagesize()` or command `getconf PAGESIZE`





# Free Frames



Before allocation

After allocation





# Implementation of Page Table

- Page table is kept in main memory
  - **Page-table base register (PTBR)** points to the **page table**
  - **Page-table length register (PTLR)** indicates **size of the page table**
- In this scheme every data/instruction access requires two memory accesses
  - One for the **page table** and **one for the data / instruction**
- The two-memory access problem can be solved by the use of a **special fast-lookup hardware** cache called **translation look-aside buffers (TLBs)** (also called **associative memory**).
  - When the associative memory is presented **with an item**, the item is compared with all keys simultaneously.





# Translation Look-Aside Buffer

---

- The TLB contains only a few of the page-table entries. When a logical address is generated by the CPU, the MMU first checks if its page number is present in the TLB.
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the **TLB for faster access next time**
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access







# Hardware

- Associative memory – parallel search

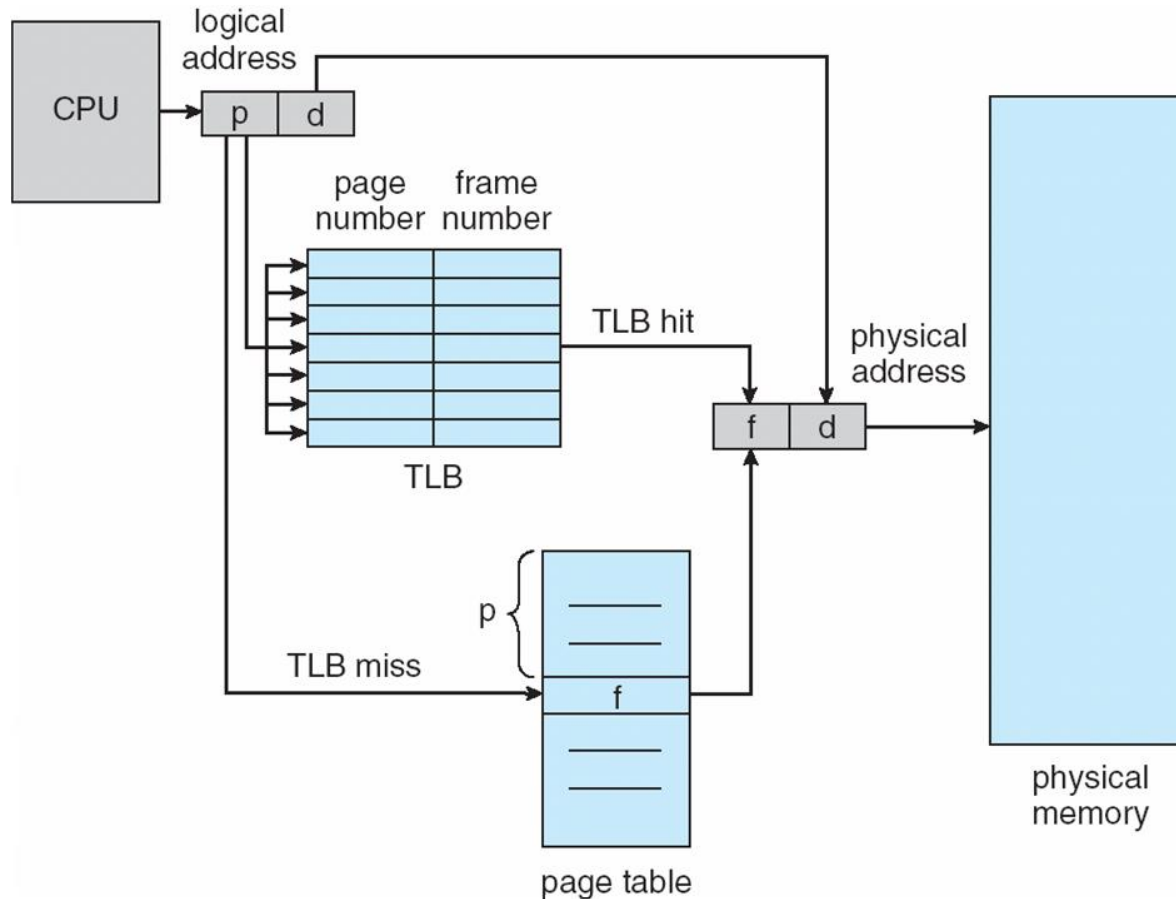
Page #	Frame #

- Address translation (p, d)
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory





# Paging Hardware With TLB





# Effective Access Time

- Hit ratio – percentage of times that a page number is found in the TLB
- An 80% hit ratio means that we find the desired page number in the TLB 80% of the time.
- Suppose that 10 nanoseconds to access memory.
  - If we find the desired page in TLB then a mapped-memory access take 10 ns
  - Otherwise we need two memory access so it is 20 ns
- **Effective Access Time (EAT)**  
$$\text{EAT} = 0.80 \times 10 + 0.20 \times 20 = 12 \text{ nanoseconds}$$

implying 20% slowdown in access time
- Consider amore realistic hit ratio of 99%,  
$$\text{EAT} = 0.99 \times 10 + 0.01 \times 20 = 10.1\text{ns}$$

implying only 1% slowdown in access time.






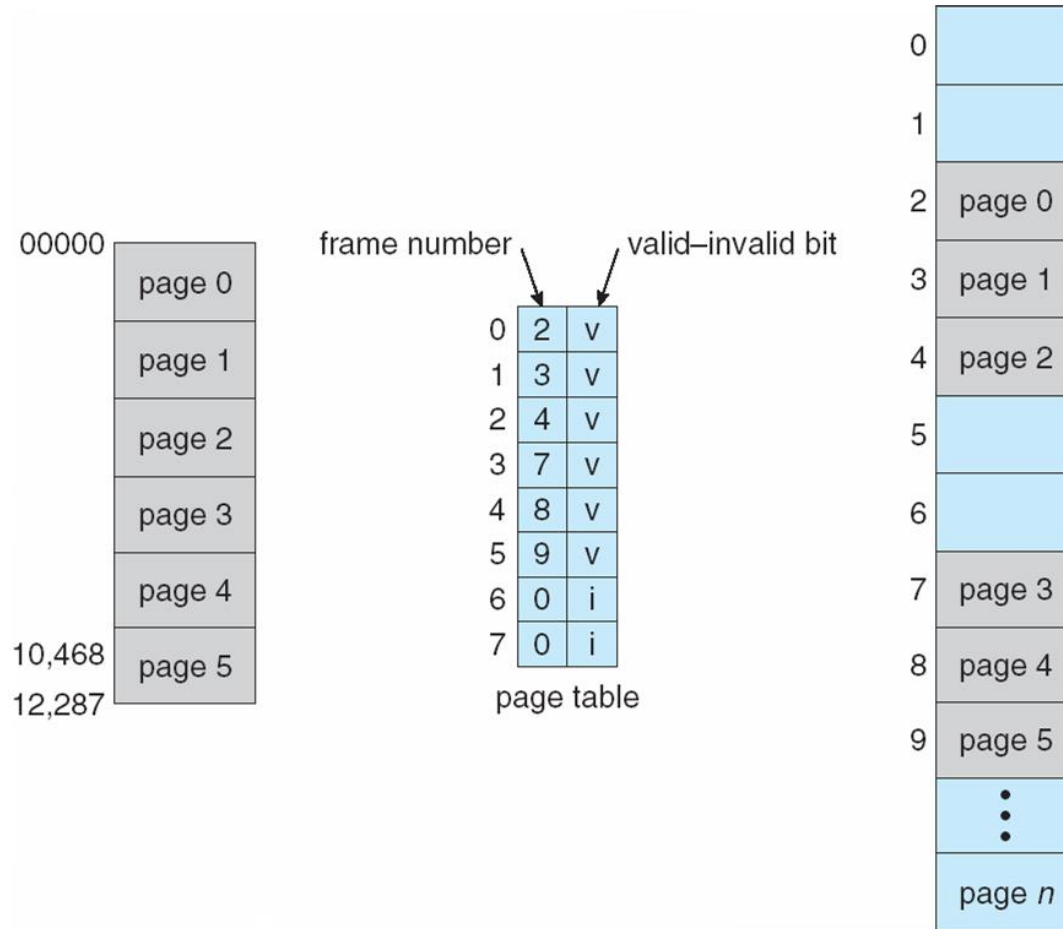
# Memory Protection

- Memory protection implemented by **associating protection bit with each frame** to indicate if **read-only or read-write access is allowed**
  - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - “invalid” indicates that the page is not in the process’ logical address space
  - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel





# Valid (v) or Invalid (i) Bit In A Page Table





# Shared Pages

---

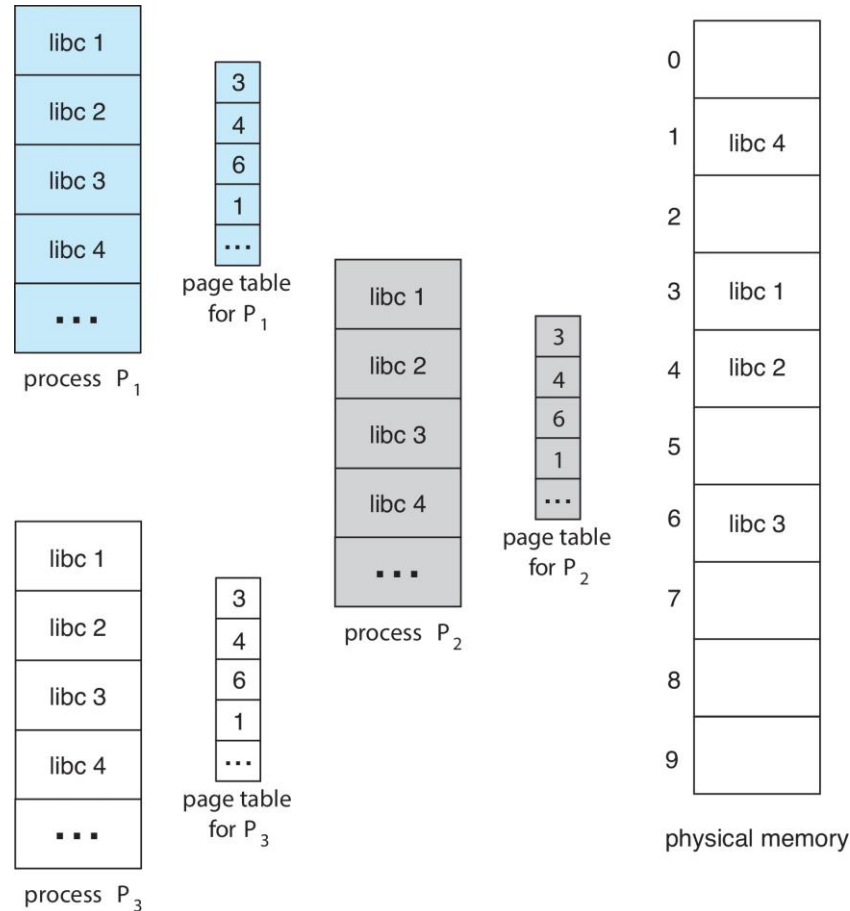
## ■ Shared code

- One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
  - ▶ Reentrant code is **non-self-modifying code**: it never changes during execution
- Similar to multiple threads sharing the **same process space**
- Also useful for **interprocess communication** if sharing of **read-write pages** is allowed





# Shared Pages Example





# Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
  - Consider a 32-bit logical address space as on modern computers
  - Page size of 4 KB ( $2^{12}$ )
  - Page table would have 1 million entries ( $2^{32} / 2^{12}$ )
  - If each entry is 4 bytes → each process 4 MB of physical address space for the page table alone
    - ▶ Don't want to allocate that contiguously in main memory
  - One simple solution is to divide the page table into smaller units
    - ▶ Hierarchical Paging
    - ▶ Hashed Page Tables
    - ▶ Inverted Page Tables

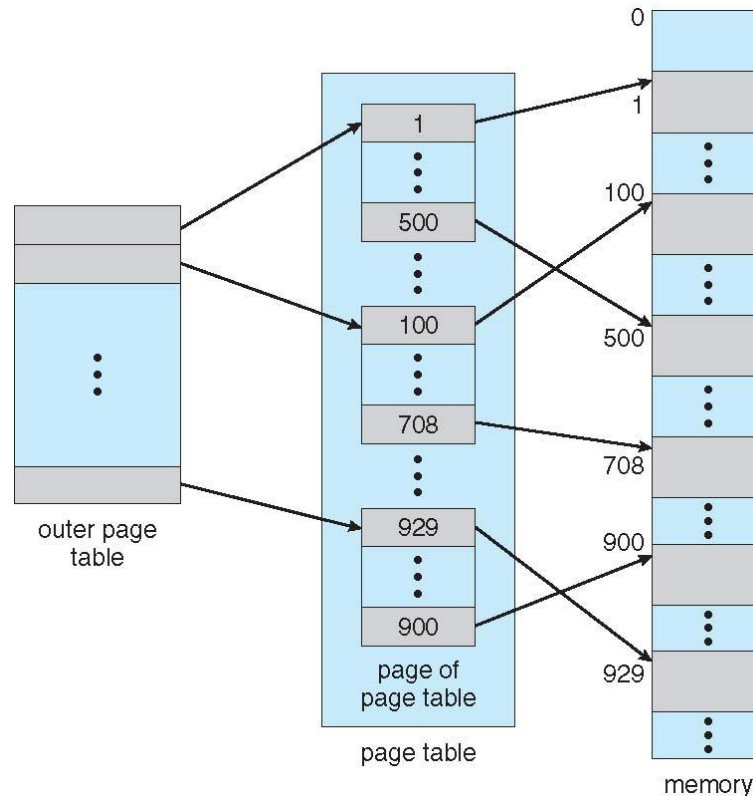






# Hierarchical Page Tables

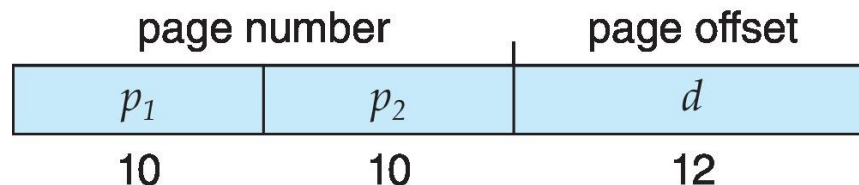
- Break up the logical address space into **multiple page tables**
- A simple technique is a **two-level page table**
- We then **page the page table**






# Two-Level Paging Example

- A logical address (on 32-bit machine with 4KB page size) is divided into:
  - a page number consisting of **20 bits**
  - a page offset consisting of **12 bits**
- Since the **page table is paged**, the page number is further divided into:
  - a 10-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:

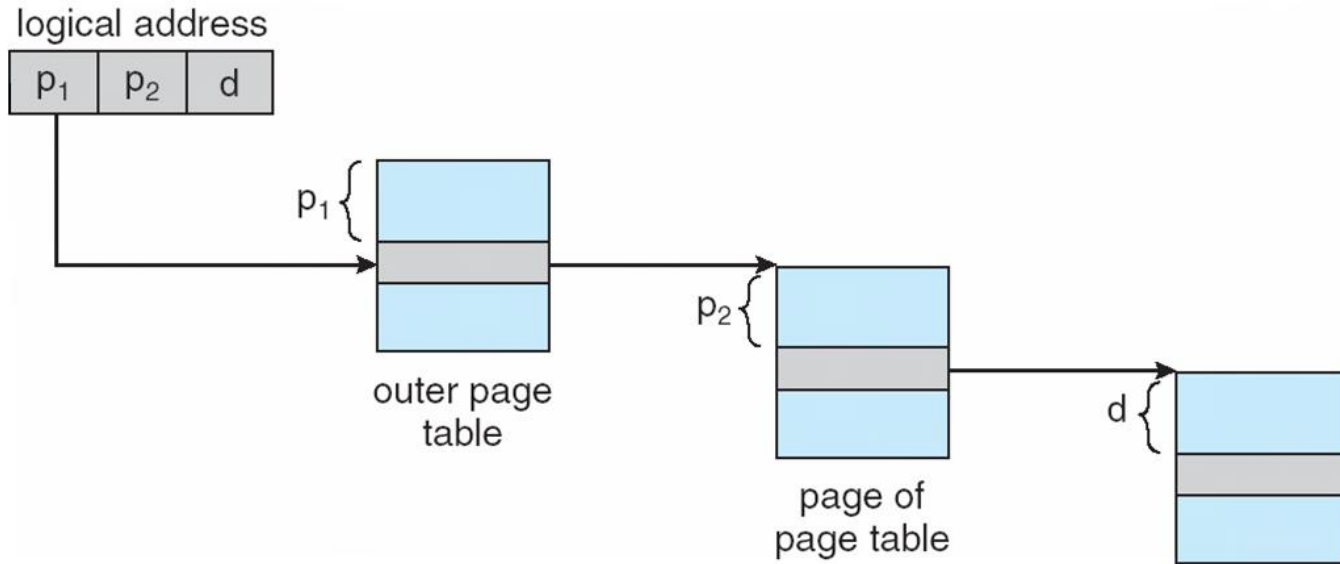


- where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the inner page table
- Known as **forward-mapped page table**





# Address-Translation Scheme





# 64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB ( $2^{12}$ )
  - Then page table has  $2^{52}$  entries
  - If two level scheme, inner page tables could be  $2^{10}$  4-byte entries
  - Address would look like

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

- Outer page table has  $2^{42}$  entries or  $2^{44}$  bytes
- One solution is to add a 2<sup>nd</sup> outer page table
- But in the following example the 2<sup>nd</sup> outer page table is still  $2^{34}$  bytes in size
  - ▶ And possibly 4 memory access to get to one physical memory location





# Three-level Paging Scheme

outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

2nd outer page	outer page	inner page	offset
$p_1$	$p_2$	$p_3$	$d$
32	10	10	12





# Hashed Page Tables


---

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
  - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
  - If a match is found, the corresponding physical frame is extracted





# Hashed Page Table





# Inverted Page Table

---

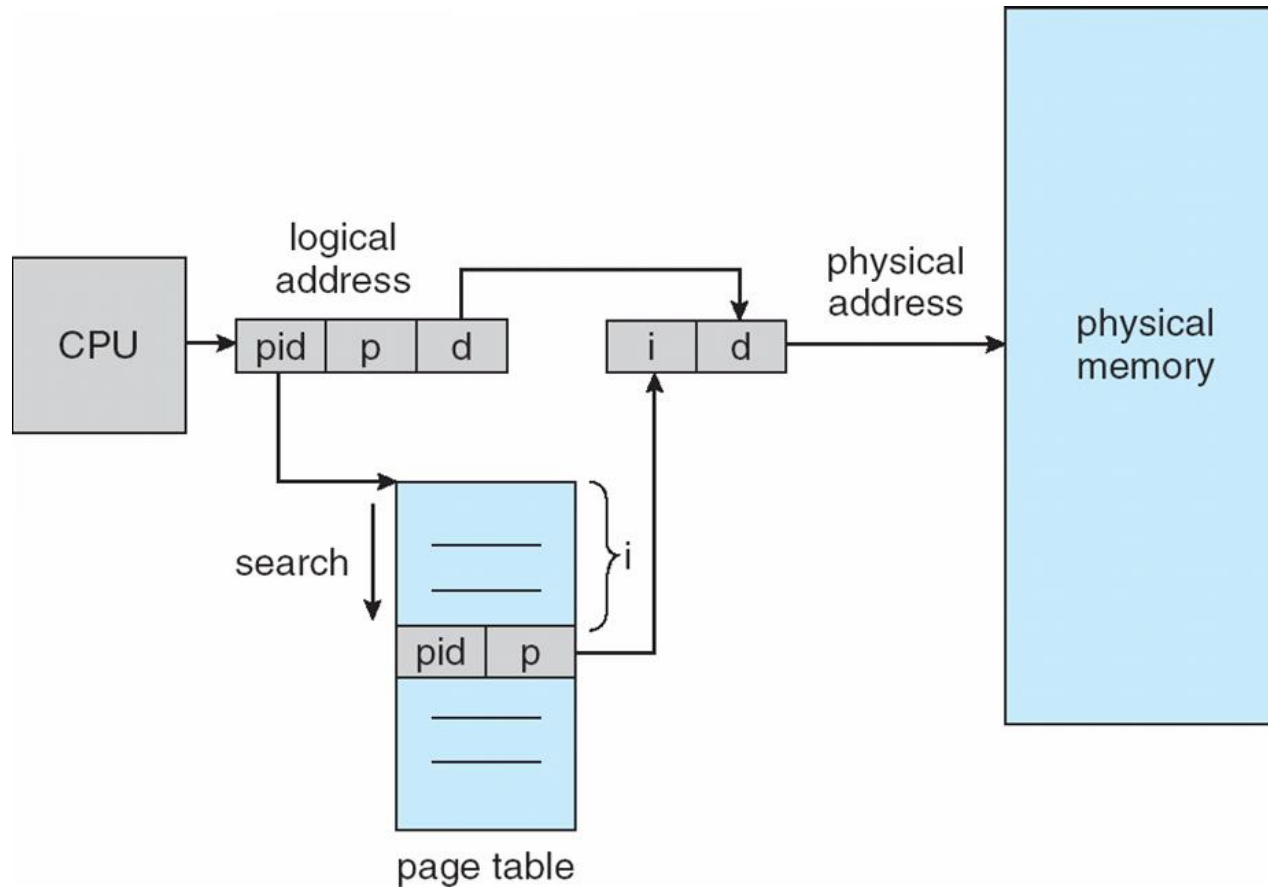
- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries
  - TLB can accelerate access
- But how to implement shared memory?
  - One mapping of a virtual address to the shared physical address







# Inverted Page Table Architecture





# Oracle SPARC Solaris

- Consider modern, 64-bit operating system example with tightly integrated HW
  - Goals are efficiency, **low overhead**
- Based on complex hashing
- Two hash tables
  - **One kernel** and one for **all user processes**
  - Each maps **memory addresses** from **virtual to physical memory**
  - Each entry represents a ***contiguous area of mapped virtual memory***,
    - ▶ More efficient than having a separate **hash-table entry for each page**
  - Each entry **has base address and span** (indicating the number of pages the entry represents)





# Oracle SPARC Solaris (Cont.)

- TLB holds translation table entries (TTEs) for fast hardware lookups
  - A cache of TTEs reside in a **translation storage buffer (TSB)**
    - ▶ Includes an entry **per recently accessed page**
- Virtual address reference causes TLB search
  - If miss, **hardware walks the in-memory TSB** looking for the TTE corresponding to the address
    - ▶ If match found, the **CPU copies the TSB entry into the TLB** and translation completes
    - ▶ If no match found, **kernel interrupted to search the hash table**
      - The kernel then **creates a TTE from the appropriate hash table** and **stores it in the TSB**, Interrupt handler returns control to the MMU, which completes the address translation by loading the TTE into the TLB.






# Swapping

- A process (or a portion of a process) can be **swapped** temporarily out of memory to a backing store, and then brought **back** into memory for continued execution
  - Total physical **memory space of processes** can exceed **physical memory**
- **Backing store** – fast disk large enough to **accommodate copies of all memory images for all users**; must provide direct access to these memory images
- **Roll out, roll in** – **swapping variant used for priority-based scheduling algorithms**; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is **transfer time**; total transfer time is **directly proportional** to the amount of memory swapped
- System maintains a **ready queue** of ready-to-run processes **which have memory images on disk**





# Schematic View of Swapping



- Does the swapped out process need to swap back in to same physical addresses?
- What if I/O is pending for the process to be swapped out?





# Standard Swapping

---

- Moving **entire processes** between main memory and a backing store.
- When a process is swapped to the backing store
  - the data structures associated with the **process must be written to the backing store.**
  - For a multithreaded process
    - ▶ per-thread **data structures** must be swapped
- The OS must also maintain **metadata for processes** that have been swapped out.
- **Advantage:** allows physical memory to be oversubscribed
- Mostly idle processes are **good candidates for swapping**
- **Disadvantage:** Transfer Time





# Swapping with Paging


---

- Standard swapping was used in **traditional UNIX systems**
  - No longer used in contemporary OSs except Solaris
    - ▶ Solaris uses standard swapping in case of **extremely low memory**
- Most contemporary OS Including Linux and Windows uses swapping with pages.
  - small number of pages will be involved in swapping.
- **page out** operation moves a page from memory to the backing store; the reverse process is known as a **page in**.





# Schematic View of Swapping with Paging







# Swapping on Mobile Systems


- Not typically supported
  - Flash memory based
    - ▶ Small amount of space
    - ▶ Limited number of write cycles
    - ▶ Poor throughput between flash memory and CPU on mobile platform
- Instead use other methods to free memory if low
  - iOS **asks** apps to voluntarily relinquish allocated memory
    - ▶ Read-only data thrown out and reloaded from flash if needed
    - ▶ Failure to free can result in termination
  - Android terminates apps if low free memory, but first writes **application state** to flash for fast restart
  - Both OSes support paging





# Example: The Intel IA-32 Architecture


- Supports both **segmentation** and **segmentation with paging**
  - Each segment can be 4 GB
  - Up to 16 K segments per process
  - Divided into two partitions
    - ▶ First partition of up to 8 K segments are private to process (kept in **local descriptor table (LDT)**)
    - ▶ Second partition of up to 8K segments shared among all processes (kept in **global descriptor table (GDT)**)






# Example: The Intel IA-32 Architecture (Cont.)

- CPU generates logical address
  - Selector given to segmentation unit
    - ▶ Which produces **linear addresses**





- Linear address given to paging unit
  - ▶ Which generates **physical address in main memory**
  - ▶ Paging units form **equivalent of MMU**
  - ▶ Pages sizes can be 4 KB or 4 MB






# Logical to Physical Address Translation in IA-32






# Intel IA-32 Segmentation






# Intel IA-32 Paging Architecture





# Intel x86-64

- Current generation Intel x86 architecture
- 64 bits is **ginormous** (> 16 exabytes)
- In practice only implement 48 bit addressing
  - Page sizes of 4 KB, 2 MB, 1 GB
  - Four levels of paging hierarchy





# Example: ARM Architecture

---

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- **Self-Study.**





# End of Chapter 9

---

